

When Agents Remember Too Much: Memory Poisoning Attacks on Large Language Model Agents

George Torres[†], Sharad Shrestha^{*}, and Satyajayant Misra[§]

^{†*§}Department of Computer Science, [§]Klipsch School of Electrical and Computer Engineering
New Mexico State University, Las Cruces, New Mexico, USA
Email: [†]gtorrez@nmsu.edu, ^{*}sharad@nmsu.edu, [§]misra@nmsu.edu

Abstract—Personal AI agents powered by large language models can reason and act using available tools to access emails, manage calendars, and push code to remote repositories, all with minimal oversight. When augmented with long-term memory, an agent can recall specific details relevant to the current task, reducing the need for large context windows. Currently, long-term memory agents tend to fall into two distinct domains: conversational and action-planning agents. Personal assistant agents sit at the convergence of these two domains and handle sensitive information while interacting with untrusted information sources, creating previously unaccounted security vulnerabilities. In this work, we introduce the novel attack vector, GhostWriter, which exploits current memory subsystems in tool-using personal agents to poison their memory store. GhostWriter operates in two phases: injection, where an adversary sends a hidden attack payload to the target agent; and activation, in which the poisoned memory is retrieved. We show that GhostWriter achieves near-universal injection rates of approximately 98% and a high average activation rate of approximately 60% against state-of-the-art agents. This attack is possible due to the lack of security-focused memory governance. In response, we propose Agentic Memory Sentry (AM-Sentry), which leverages two mitigation techniques: a memory-saving policy and a memory-retrieval screen. Our experiments show that AM-Sentry dramatically reduces GhostWriter’s success rate while preserving agent utility.

1. Introduction

Large language model (LLM) agents can reason and act with minimal human oversight, enabling automation of complex, multi-step tasks. These agents have access to tools that improve their efficiency and agency, allowing them to access emails, reply to calendar invites, and even push code to remote repositories [1], [2]. An agent’s reasoning, decision-making, and accuracy can be further improved through the integration of long-term memory [3], [4]. An agent with long-term memory can recall specific details relevant to its current task, reducing the need for large context windows that do not scale well. Figure 1 illustrates this with an example: across sessions or after long periods between related tasks, an agent with only its context window




Figure 1: Stateless LLM agents process each interaction in isolation. LLM agents with persistent memory allow for querying and retrieving saved information.

struggles to accurately answer the user’s question and may even hallucinate. In contrast, the agent equipped with a long-term memory store can dynamically retrieve the necessary context and accurately answer the given question.

The existing literature on agents with long-term memory has primarily been split into two distinct domains: conversational agents (i.e., chatbots) [3], [5], whose primary role is to talk with a user, and action-planning agents [4], [6], which focus on improving tool use or other complex operations, such as robotic control or code execution. Practical use cases that combine aspects of both domains, such as personal assistant agents or automated coding agents [7], [8], are currently underrepresented, and these agents would clearly benefit from including both fact-based and action-planning memory for tool use. When long-term memory is used in such agents, new security risks arise in the context of untrusted sources of information, their storage, and how they may, in turn, affect the agent’s future behavior.

These agents routinely interact with untrusted sources of information, such as emails and documents, creating a distinct attack surface for indirect memory poisoning. An adversary can send malicious or misleading content to an agent (e.g. email agent), which is then stored in the agent’s memory store for later use. Later, when the user submits

当智能体记忆过载时：针对大型语言模型智能体的内存中毒攻击

乔治·托雷斯[#]、沙拉德·什雷斯塔^{*}与萨蒂亚贾扬特·米斯拉[§]

[#]^{*}§ 计算机科学系, [§]Klipsch 电气与计算机工程学院

美国新墨西哥州拉斯克鲁塞斯, 新墨西哥州立大学

电子邮件: [#]gtorrez@nmsu.edu, ^{*}sharad@nmsu.edu, [§]misra@nmsu.edu

摘要——基于大型语言模型驱动的个人人工智能代理能够利用可用工具进行推理与操作, 例如访问电子邮件、管理日历以及将代码推送到远程代码库, 且整个过程几乎无需人工监督。当引入长期记忆功能后, 该代理可召回与当前任务相关的具体细节, 从而降低对宽广上下文窗口的需求。目前, 基于长期记忆的智能代理通常分为两个 **distinct** 领域: 对话型代理与行动规划型代理。个人助理代理处于这两个领域之间的交汇点, 在与不可信的信息源交互时处理敏感信息, 从而引发此前未被充分考虑的安全漏洞。在本研究中, 我们提出一种名为 **GhostWriter** 的新型攻击向量, 该攻击向量利用工具使用型个人智能代理中的当前内存子系统来污染其内存存储空间。**GhostWriter** 的攻击过程分为两个阶段: 注入阶段——攻击者向目标代理发送隐藏的攻击载荷; 以及激活阶段——从受污染的内存中提取攻击载荷。我们证明, 与当前最先进的智能体相比, **GhostWriter** 实现了接近普遍性的注入率 (约 98%) 以及较高的平均激活率 (约 60%)。这一攻击之所以可行, 是因为当前系统缺乏针对安全性的内存管理机制。作为应对措施, 我们提出了“智能体内存哨兵” (**AM-Sentry**), 该方案采用了两种缓解技术: 一种是内存节省策略, 另一种是内存检索筛选机制。实验结果表明, **AM-Sentry** 能显著降低 **GhostWriter** 的攻击成功率, 同时仍能保障智能体的效用。

1. 引言

大型语言模型 (LLM) 智能体可在极少人工干预的情况下进行推理与操作, 从而实现复杂、多步骤任务的自动化。这些智能体可调用各类工具以提升其效率与自主性, 使其能够查看电子邮件、回复日历邀请, 甚至将代码推送到远程代码库[1]。[2]通过整合长时记忆[3], 可进一步提升智能体的推理能力、决策能力及准确性。[4]具备长期记忆能力的智能体能够召回与其当前任务相关的具体细节, 从而减少了对那些扩展性较差的大型上下文窗口的依赖。图1通过一个示例阐明了这一点: 在跨会话使用场景或在相关任务之间经过较长时间间隔




图 1: 无状态大语言模型 (LLM) 智能体在每次交互中均独立处理; 而具备持久化记忆功能的 LLM 智能体则支持查询及检索已保存的信息。

后, 仅依赖于自身上下文窗口的智能体往往难以准确回答用户的提问, 甚至可能出现幻觉现象; 相比之下, 配备长期记忆存储机制的智能体能够动态检索所需上下文信息, 并准确回答给定问题。

现有关于具有长期记忆的智能体的文献主要可分为两个 **distinct** 领域: 会话智能体 (即聊天机器人) [3], [5] 其主要职责是与用户进行对话, 以及执行规划智能体 [4], [6] 这些系统侧重于优化工具使用或其他复杂操作 (例如机器人控制或代码执行)。而结合了这两个领域特性的实际应用案例 (例如个人助理智能体或自动化编程智能体[7]), [8] 这些智能体目前在系统中的代表性不足; 若能在其工具使用机制中同时引入基于事实的记忆与行动规划记忆, 这些智能体显然将获益匪浅。当此类智能体调用长期记忆时, 在面对不可信的信息源、这些信息的存储方式以及这些因素可能对智能体未来行为产生的影响时, 便会产生新的安全风险。

这些智能体通常会与不可信的信息源 (如电子邮件或文档) 进行交互, 从而为间接内存中毒攻击创造了独特的攻击面。攻击者可向智能体 (例如电子邮件智能体) 发送恶意或具有误导性的内容, 这些内容随后会被存储在该智能体的内存存储区中以供后续使用。此后,

a benign prompt, the malicious content is retrieved from the memory store and added to the agent’s context, which alters the agent’s behavior to align with the adversary’s goals. For example, consider a personal assistant agent that saved a malicious workflow instruction that later directs it to include an unauthorized recipient in emails containing confidential information about the user’s current project, thereby compromising the user’s nondisclosure agreements.

Such attacks are enabled by the lack of security-focused memory governance. Currently, long-term memory agents tend to retain every interaction. Memory governance for agents has been proposed by A-MAC [9], though it primarily focuses on improving utility rather than detecting malicious or misleading content, leaving the underlying memory stores vulnerable to these attacks. Security-focused memory governance is needed to protect agents from the influence of malicious content while remaining productive.

The poisoning of an agent’s memory store or knowledge base has been explored by works such as AgentPoison [10] and MINJA [11], which demonstrated that corrupted memory can significantly degrade agent behavior and can be exploited to induce harmful actions. The adversaries in these works have direct access to the memory system or direct interaction with the target agents, allowing them to directly inject memories or prompt agents to store poisoned memories. A more realistic scenario assumes that an adversary would not have direct access to the agent or its memory store; instead, they must poison the agent’s memory through the inputs it receives, such as emails, documents, and calendar events.

To address these gaps in the literature, we detail a novel attack vector *GhostWriter*, which exploits the lack of security-focused memory governance in tool-using personal assistant agents to poison the memory store. *GhostWriter* operates in two phases: injection, where an adversary sends a hidden attack payload to the target agent, followed by activation, in which the poisoned memory is retrieved and influences the agent’s behavior. We show that *GhostWriter* achieves average injection rates of roughly 98% and a high average activation rate of roughly 60% across all tested state-of-the-art agent architectures. We also compare *GhostWriter* to prompt injection [12], another type of attack against agents, and highlight the differences in attack signature and activation, showing that *GhostWriter* has a greater impact on the evaluated agents.

We also propose Agentic Memory Sentry (*AM-Sentry*), which combines two mitigation techniques: a memory-saving policy and a retrieval screen, designed to prevent these attacks while maintaining utility. *AM-Sentry*’s memory-saving policy has three levels of strictness, each adding stronger constraints and vulnerability-detection capabilities. The retrieval screen can be enabled or disabled. When enabled, it reviews each retrieved memory entry and filters out those that contain vulnerabilities or hidden directives that could alter the agent’s behavior. The three memory-saving policies, paired with the two retrieval screen settings, yields six configurations. We analyze the effectiveness of each configuration in preventing *GhostWriter* attacks

and its effects on agent performance. We show that attack effectiveness can be dramatically reduced at minimal cost to agent utility. We also introduce a custom test suite that assesses an agent’s recall and task performance in a personal assistant use case, allowing us to assess the impact *AM-Sentry* has on each agent’s base utility.

Our **novel contributions** include:

- We detail *GhostWriter*, a novel attack vector that exploits the lack of security-focused memory governance in long-term memory tool-using agents to poison the memory store via untrusted tool inputs.
- We propose *AM-Sentry*, a novel memory defense framework with six configurations of memory policy strictness and retrieval screening.
- We evaluate *GhostWriter* against five state-of-the-art agents across four LLM models and assess the effectiveness of *AM-Sentry* in reducing attack success and preserving agent utility.

We plan to release our implementation and evaluation toolkit, which includes our attack setup and agent performance test suite, to spur others to test new attack vectors and mitigation techniques or to evaluate the utility of their agent implementation.

The rest of the paper is organized as follows. Section 2 provides a literature review on LLM agents and memory poisoning, while Section 3 defines our system and threat models. We present the *GhostWriter* attack vector in Section 4 and the *AM-Sentry* mitigation architecture in Section 5. Our experimental methodology and evaluations are then covered in Sections 6 and 7, respectively. Finally, Sections 8 and 9 detail limitations, future work, and our concluding remarks.

2. Background and Related Work

2.1. LLM Agents

LLM agents are autonomous systems that leverage LLMs to perceive, reason, and take dynamic actions based on their current context [2]. Agents are increasingly deployed in real-world settings [13], with open-source frameworks such as AutoGPT [14] and BabyAGI [15] providing easy means of adoption. Multi-step reasoning agents, such as ReAct [16], interleave reasoning with actions, enabling them to generate thoughts that enrich their current context and improve later decisions. Equipping agents with external tools provides them with a structured means of interacting with external elements and improving their performance [1]. Tools are widely used and incorporated into production-level LLM interfaces such as ChatGPT and Claude [7], [17]. Combining tool use and multi-step reasoning has greatly increased agent utility across the board [2], [18].

2.2. Memory-based Agents

Memory context is a prevailing issue in current agents and LLM interfaces; large context windows tend to slow

当用户提交一个良性提示时，系统会从内存存储区中检索出该恶意内容并将其添加到智能体的上下文中，进而改变智能体的行为以符合攻击者的目标。例如，假设一个个人助理智能体保存了恶意工作流指令，该指令随后指示该智能体在包含用户当前项目机密信息的电子邮件中添加未经授权的收件人，从而导致用户违反其保密协议。

此类攻击源于缺乏以安全为导向的记忆管理机制。目前，长时记忆智能体往往会保留所有交互记录。AM-MAC[9]曾提出针对智能体的记忆管理方案，但该方案主要侧重于提升效用而非检测恶意或误导性内容，从而使得底层记忆存储库容易受到此类攻击的影响。因此，亟需引入以安全为导向的记忆管理机制，以保护智能体免受恶意内容的影响，同时确保其持续高效运行。

智能体内存存储器或知识库遭中毒的问题已得到多项研究工作的探讨，例如 AgentPoison[10]与 MINJA[11]；这些研究证明，受损的内存会显著影响智能体的行为，并可被利用以诱使智能体执行有害操作。在上述研究中，攻击者能够直接访问内存系统或与目标智能体进行直接交互，从而直接注入恶意内存或诱导智能体存储受污染的内存。而在一个更为现实的场景中，攻击者无法直接访问智能体或其内存存储器；相反，他们必须通过智能体接收的输入（如电子邮件、文档及日历事件）来对智能体的内存进行“中毒”操作。

为弥补现有文献中的这些研究空白，我们详细阐述了一种新型攻击向量——*GhostWriter*，该攻击利用了工具使用型个人助理智能体中缺乏安全导向的内存管控机制这一漏洞，从而对内存存储库进行中毒攻击。*GhostWriter*的攻击过程分为两个阶段：注入阶段——攻击者向目标智能体发送隐藏的攻击载荷；随后进入激活阶段——从中毒的内存中检索数据并影响智能体的行为。我们证明，在所有经过测试的先进智能体架构上，*GhostWriter*的平均注入成功率约为 98%，平均激活成功率高达约 60%。我们还将*GhostWriter*与另一种针对智能体的攻击方式——提示词注入[12]进行对比，并重点分析了两者在攻击特征与激活效果上的差异，结果表明*GhostWriter*对所评估的智能体具有更显著的影响。

我们还提出了 *AM-Sentry*，该方案结合了两种缓解技术：内存节省策略与检索筛选机制，旨在在保障系统效用的同时防范此类攻击。*AM-Sentry*的内存节省策略设定了三个不同严格度级别，每个级别均附加了更严格的约束条件及漏洞检测能力；而检索筛选机制可选择启用或禁用。当启用该机制时，系统会对每个检索到的内存条目进行审查，并过滤掉其中包含漏洞或可能改变智能体行为的隐藏指令的条目。这三种内存节省策略与两种检索筛选机制的组合方式共可形成六种配置方案。我们分析了每种配置在防范*GhostWriter*攻击方面的有效性及

其对智能体性能的影响。结果表明，在将智能体效用损失降至最低的前提下，可显著降低攻击的有效性。此外，我们还引入了一套定制化测试套件，用于评估智能体在个人助理应用场景中的召回率与任务执行性能，从而能够量化*AM-Sentry*对各智能体基础效用的影响。

我们的新贡献包括：

- 我们详细介绍了 *GhostWriter*——一种新型攻击向量，该向量利用了长期使用内存工具的智能体在内存管理方面缺乏安全防护机制这一漏洞，通过利用不可信的工具输入来对内存存储库进行中毒攻击。
- 我们提出 *AM-Sentry*——一个新型内存防御框架，该框架支持六种不同内存策略严格度与检索筛选的配置方案。
- 我们针对*GhostWriter*与五款最先进的智能体，在四种大语言模型（LLM）上进行对比评估，并检验 *AM-Sentry* 在降低攻击成功率及维护智能体效用方面的有效性。

我们计划发布我们的实现与评估工具包（其中包含攻击场景搭建工具及智能体性能测试套件），旨在激励其他研究者测试新的攻击向量与缓解技术，或评估其智能体实现方案的实用性。

本文其余部分结构安排如下：第2节对LLM智能体与内存中毒问题进行文献综述；第3节定义了本文所提出的系统及威胁模型；第4节介绍*GhostWriter*攻击向量；第5节阐述*AM-Sentry*缓解架构；随后，第6节与第7节分别阐述实验方法与评估工作；最后，第8节与第9节详细说明本文的局限性、未来工作方向及结论。

2. 背景与相关工作

2.1. LLM 代理

LLM智能体是基于大型语言模型（LLM）进行感知、推理并根据当前上下文采取动态行动的自主系统[2]。这类智能体正日益广泛地部署于现实场景中[13]，而诸如AutoGPT[14]和BabyAGI[15]等开源框架为其部署提供了便捷途径。多步推理型智能体（如ReAct[16]）将推理过程与具体行动交替执行，使其能够生成有助于丰富当前上下文并优化后续决策的思考过程。为智能体配备外部工具为其提供了与外部元素交互并提升性能的结构化手段[1]。此类工具已被广泛采用，并集成到ChatGPT和Claude[7]等生产级LLM接口中。[17]将工具使用与多步骤推理相结合，显著提升了智能体在各领域的效用[2]。[18]。

2.2. 基于内存的智能体

内存上下文是当前智能体及大型语言模型（LLM）接

down performance and increase computational costs [3]. To address this issue, systems that manage an agent’s current context via an additional database have been proposed. One such approach is the Retrieval-Augmented Generation (RAG) system, which pulls in real-time context from a pre-seeded database [19] based on the given prompt. A variety of RAG systems have been proposed [20] in the literature, but they mostly rely on pre-existing information and are unable to learn from their experiences.

RAG systems and other information stores benefit from agents saving their own examples, experiences, and acquired knowledge in them. The concept of long-term cross-session memory was introduced by works such as MemGPT [3] and Generative Agents [21], showing that improvements in agent response and recall greatly outweigh those of simple systems that rely on longer context windows, which are not scalable. Agents store knowledge/interactions in their memory stores, labeling them with descriptive tags that enable later retrieval based on semantically similar prompts. MemoryOS [5] introduced a memory system composed of three types of memory stores: short-, medium-, and long-term. A-MEM [22] introduces dynamic memory linking, which uses an LLM call to help manage the memory store and the relationship between memories.

The works discussed thus far have focused on conversational agents that engaged in conversation with a single user; a separate line of work involves action-planning agents such as Reflexion [23], which uses repeated self-reflections on past attempts to improve its current action plan. These action plans represent the steps an agent takes to reach a goal, which could involve successfully using an external tool or the steps needed to move a robotic machine. Retrieval-augmented planning (RAP) [6] introduces a long-term memory store for these types of agents, allowing them to draw on previously successful attempts (trajectories) at the same or similar tasks to improve performance and reduce the number of attempts needed to succeed. ExpeL [4] pairs both successful and failed stored trajectories with reflections to further improve performance, allowing agents to extract insights from their past experiences.

Current work on long-term memory agents can be separated into two domains: conversational agents, which remember facts about their users, and action-planning agents, which remember optimal means for performing various actions and tool calls. Both domains focus on improving how memories are stored, tagged, and retrieved within their own context, leaving their convergence underexplored. Realistic, practical deployment, such as an automated coding or a personal assistant agent, is missing from these works. These types of agents require aspects from both domains, as they must store information about their users or work environments and use tools to perform specific actions.

2.3. Memory Poisoning Attacks and Defenses

Personal assistant agents are responsible for handling sensitive information and operations; thus, it is critical that their behavior not be steered toward adversarial outcomes.

Agents powered by LLMs have been shown to be easily manipulated when provided with false or misleading information, as demonstrated by recent incidents involving Meta customer support chat agents [24]. AgentPoison [10] shows the potential damage when such false or misleading information is persisted into long-term memory, showing that an injected poisoned memory can hijack an agent’s actions and induce malicious behavior. Agents can also be tricked into executing and then saving a malicious action plan through careful prompt injection, as detailed by MINJA [11]. These attacks center on impractical adversarial models; either the adversary has direct access to the memory store or to the agent itself, which would require access to the user’s credentials to mount an attack.

Another means of mounting a memory poisoning attack is to feed an agent corrupted inputs through its tool interface, which are saved as part of an agent’s regular operations. Some early work has considered this attack vector, such as Zombie Agents [25], which details how web search-capable agents exposed to malicious web pages can be manipulated into storing adversarial instructions that alter behavior upon retrieval. However, this work operates under assumptions that limit its applicability to realistic deployments; it requires an agent with RAG-based memory to repeatedly store the same attack payload hundreds of times to be effective, something that would not work on memory systems with simple deduplication. This and similar works, such as MemoryGraft [26], focused their attack evaluations on their own agent implementations rather than demonstrating their applicability to the state-of-the-art.

Defenses against these poisoning attacks are an emerging area with limited dedicated work. The existing literature tends to focus on AgentPoison and MINJA-type attacks. A-MemGuard [27], for example, proposes a consensus-based retrieval mechanism to detect anomalous memories, an approach that is inapplicable to fact-based memory systems where no comparable reasoning paths exist. A-MAC [9] introduces the concept of memory governance, but focuses solely on utility and does not address memory poisoning vulnerabilities. No existing work investigates indirect memory poisoning through untrusted tool inputs against state-of-the-art agents, nor proposes defenses that address this attack path or evaluate the resulting security-utility tradeoff. In this paper, we attempt to plug this gap.

3. Models and Assumptions

In this section, we describe the models and assumptions underlying our work on personal AI agents. We first present the system model, then outline our security assumptions, and finally define the threat model.

3.1. System Model

We envision a practical deployment in an office environment where a personal AI agent assists with email and calendar management. The agent autonomously monitors the user’s inbox and activates when the user receives emails or

口中一个普遍存在的问题：过大的上下文窗口往往会降低系统性能并增加计算成本[3]。为解决这一问题，已有研究提出通过额外数据库来管理智能体当前上下文的系统。其中一种方法是检索增强生成（RAG）系统，该系统根据给定的提示词，从预置数据库中实时提取上下文信息[19]。目前，文献中已提出了多种 RAG 系统[20]，但这些系统大多依赖于预先存在的信息，无法从自身经验中进行学习。

RAG 系统及其他信息存储库能够受益于智能体将自身的示例、经验及习得的知识存入其中。长期跨会话记忆这一概念由 MemGPT[3]和 Generative Agents[21]等研究工作提出，这些研究表明，相较于依赖于不可扩展的较长上下文窗口的简单系统，智能体在响应和检索方面的性能提升要显著更大。智能体将知识或交互记录存储在其记忆库中，并为其标注描述性标签，以便基于语义相似的提示进行后续检索。MemoryOS[5]引入了一种由短期、中期和长期三种类型记忆库组成的记忆系统。A-MEM[22]提出了动态记忆链接机制，该机制利用大语言模型（LLM）调用来辅助管理记忆库以及记忆之间的关联关系。

迄今为止讨论的各类研究均聚焦于与单个用户进行对话的对话式智能体；另一条研究路径则涉及动作规划智能体，例如 Reflexion[23]——该智能体通过对过往尝试进行反复自我反思来优化当前的行动方案。这些行动方案代表了智能体为实现目标所采取的步骤，可能包括成功使用外部工具或执行机器人机械臂运动所需的步骤。检索增强规划（RAP）[6]为这类智能体引入了长期记忆存储机制，使其能够借鉴在相同或类似任务中取得成功的历史尝试（轨迹），从而提升性能并减少达成目标所需的尝试次数。Expel[4]将成功与失败的存储轨迹与反思过程相结合，以进一步优化性能，使智能体能够从过往经验中汲取启示。

当前针对长期记忆智能体的研究可划分为两个领域：一是对话智能体——这类智能体能够记住有关用户的信息；二是动作规划智能体——这类智能体能够记住执行各类操作及调用工具的最优方案。这两个领域均致力于优化记忆在各自应用场景中的存储、标记与检索机制，但二者之间的融合研究尚不充分。目前这些研究尚未涵盖实际场景下的应用部署，例如自动化编程系统或个人助理智能体等。此类智能体需要融合两个领域的技术特性——它们既需存储有关用户或工作环境的信息，又需利用工具来执行特定操作。

2.3. 内存中毒攻击与防御

个人助理智能体负责处理敏感信息及各类操作；因此，确保其行为不被引导至对抗性结果至关重要。研究表明，当智能体接收到虚假或误导性信息时，极易被操

纵——这一点在近期涉及 Meta 客户支持聊天智能体的事件中得到了印证[24]。AgentPoison[10]展示了当此类虚假或误导性信息被持久存储于智能体的长期记忆中时可能造成的潜在危害，表明注入的中毒记忆可劫持智能体的行为并诱发恶意行为。此外，通过精心设计的提示词注入技术，智能体亦可能被诱使执行并随后保存恶意行动方案，如 MINJA[11]所详述。这些攻击手段均基于不切实际的对抗性模型；此时，攻击者要么直接访问内存存储库，要么直接操控智能体本身，而后者需要获取用户凭证才能发起攻击。

另一种实施内存中毒攻击的手段是通过智能体的工具接口输入受污染的输入数据——这些输入数据通常作为智能体常规操作的一部分被保存。部分早期研究已探讨过这一攻击向量，例如《Zombie Agents》[25]，该研究详细阐述了：当具备网页搜索能力的智能体遭遇恶意网页时，如何被诱使存储对抗性指令，从而在数据检索时改变其行为。然而，这项研究基于某些假设，限制了其在实际部署场景中的适用性；它要求智能体配备基于 RAG 的内存系统，并需将相同的攻击载荷重复存储数百次才能奏效——而在采用简单去重机制的内存系统中，这一方法无法实现。上述研究以及类似的研究（如 MemoryGraft[26]）在进行攻击效果评估时，主要针对其自身实现的智能体展开，而非验证该攻击方法对当前最先进技术系统的适用性。

针对此类内存中毒攻击的防御机制是一个新兴领域，但目前针对该领域的专门研究尚不充分。现有文献多集中于 AgentPoison 和 MINJA 类攻击；例如，AMemGuard[27]提出了一种基于共识的检索机制用于检测异常内存，但该方法不适用于缺乏可比推理路径的事实型内存系统。A-MAC[9]引入了“内存治理”概念，但其研究仅侧重于效用性，未涉及内存中毒漏洞的防御问题。目前尚无研究探讨通过不可信的工具输入实施的间接内存中毒攻击对先进智能体的影响，亦未提出针对此类攻击路径的防御方案或对由此产生的安全与效用之间的权衡关系进行评估。本文旨在填补这一研究空白。

3. 模型与假设

本节阐述了我们针对个人人工智能智能体开展研究所依据的模型与假设。我们首先介绍系统模型，随后概述我们的安全假设，最后定义威胁模型。

3.1. 系统模型

我们设想该系统可在办公环境中实际部署，由个人 AI 代理协助用户管理电子邮件和日历事务。该代理可自主监控用户的收件箱，并在用户收到电子邮件或日历邀请时自动激活。用户还可直接向该代理发送指令，例如




Figure 2: AI email assistant with persistent memory.

calendar invites. The user can also submit prompts directly to the agent to request email summaries, feedback on in-progress documents, reminders of upcoming meetings, or to send emails on their behalf. To support interactions with external services, the agent has access to tool APIs that enable email and calendar functionality. The agent is also equipped with a long-term memory store to improve recall and better support task continuity and context retention.

We consider four entities in this deployment: the user, their agent, internal contacts, and external contacts. Agents remain dormant until an event triggers their response. Events include direct user input, incoming emails or calendar invites, and scheduled events. When triggered, an agent processes the interaction and writes it to memory as a summary, extracted facts, a useful trajectory, or verbatim text, depending on the agent’s design. During processing, agents tag each memory based on its content and link it to closely related memories.

When the user later submits a task or question, the agent retrieves relevant information into its context and responds accordingly. Retrieval typically proceeds by matching the tags in the user’s query to those in the memories. Figure 2 illustrates a use case for an AI email assistant: the agent retrieves the relevant meeting invite, recalls a deadline change, and drafts an email to the correct recipient. An agent without memory lacks this context and would fail to recall the deadline change or identify the correct recipient.

3.2. Security Assumptions

Our security analysis is predicated on a set of operational assumptions. First, we assume the agent’s core execution logic is benign; its hardcoded actions operate in good faith and are free of malicious code. Second, we assume the agent’s memory infrastructure maintains strict cryptographic or logical isolation, rendering it fully resilient to out-of-band tampering thus functioning as a black box to the adversary. Finally, we assume user accounts are secure, and external adversaries cannot bypass standard access controls to directly interface with the agent or its host environment.

3.3. Threat Model

We consider the user to be fully trusted, but it’s agent can be steered toward malicious actions if its context is compromised by adversarial directives. We also consider internal and external adversaries. Internal adversaries are parties with an established working relationship with the user’s organization, including employees, contractors, vendors, and other business partners. External adversaries are third parties with no such relationship. Both classes share the same attack mechanism: they use disguised inbound messages, such as emails that appear benign at a glance but contain hidden instructions. They differ, however, in contextual knowledge. Internal adversaries possess knowledge of the user’s work environment, including organizational structure, colleague names, ongoing projects, and project-specific context, making their attacks more convincing. External adversaries are limited to publicly available information, such as professional networking sites and company websites. Neither class can directly access the agent’s API, modify the memory store, or compromise the user’s account.

The adversary attempts to get the agent to store malicious or misleading information in its persistent memory, thereby influencing its future behavior. The adversary’s goals fall into four categories:

- 1) **Integrity corruption:** the adversary injects misleading information, causing the agent to produce incorrect outputs or take unintended actions in future interactions.
- 2) **Sensitive information leakage:** the adversary induces the agent to disclose the user’s private information.
- 3) **Covert exfiltration:** the adversary plants instructions that cause the agent to covertly transmit information about the user’s activities to an external party.
- 4) **Scope escalation:** the adversary causes the agent to take actions beyond its intended scope.

4. GhostWriter Architecture

In this section, we elaborate on *GhostWriter*’s mechanisms and walk through formalized protocols. *GhostWriter* follows a two-phase protocol: **injection** (Section 4.1), where the adversary feeds a poisoned memory to the target, and **activation** (Section 4.2), where the poisoned memory is retrieved on a benign query, altering the agent’s behavior. Note that the adversary controls only the injected content; they have no influence on future user prompts or on the agents’ retrieval mechanisms. Figure 3 gives an overview of both phases where an adversary sends an email claiming that a colleague’s email address has changed. This information is stored by the agent and later retrieved when the user requests to send an email to their colleague, at which point the agent redirects the email to the malicious address.

4.1. Memory Injection

GhostWriter’s first phase is formalized in Protocol 1, which begins with adversary $\mathcal{A}$ choosing a goal (g) (Line




图 2：支持持久化内存的 AI 电子邮件助手。

请求获取电子邮件摘要、对当前处理中的文档提供反馈、提醒即将召开的会议，或代为发送电子邮件。为支持与外部服务的交互，该代理可调用支持电子邮件及日历功能的工具API；此外，该代理还配备了一个长期记忆存储库，以增强信息检索能力，并更好地支持任务连续性与上下文保持。

在本次部署中，我们考虑了四个实体：用户、其智能代理、内部联系人以及外部联系人。智能代理在事件触发其响应之前均处于休眠状态；这些事件包括用户的直接输入、收到的电子邮件或日历邀请，以及已安排的事件。当事件被触发时，智能代理会处理此次交互，并根据其设计要求，将该交互记录到内存中——可作为交互摘要、提取的事实、有用轨迹或逐字文本。在处理过程中，智能代理会根据每条记忆的内容对其进行标记，并将其与相关联的记忆进行链接。

当用户随后提交任务或提出问题时，智能体会从其上下文中检索相关信息，并据此作出响应。该检索过程通常通过将用户查询中的标签与记忆库中的标签进行匹配来实现。图2展示了一个 AI 电子邮件助手的用例：该智能体会检索到相关的会议邀请函，召回一项截止日期变更信息，并为正确的收件人起草一封电子邮件。若智能体缺乏记忆功能，则无法获取此类上下文信息，因而无法召回该截止日期变更信息或识别出正确的收件人。

3.2. 安全假设

我们的安全分析基于一系列运行时假设。首先，我们假设智能体的核心执行逻辑是良性的；其硬编码操作均基于善意执行，且不含恶意代码。其次，我们假设智能体的内存基础设施实现了严格的加密或逻辑隔离，从而使其完全抵御带外篡改行为，对攻击者而言相当于一个“黑盒”。最后，我们假设用户账户安全可靠，外部攻击者无法绕过标准访问控制机制而直接与智能体或其宿主环境进行交互。

3.3. 威胁模型

我们假设用户是完全可信的，但若其上下文环境受到对抗性指令的影响，其代理程序仍可能被引导执行恶意操作。我们同时考虑内部与外部攻击者：内部攻击者是指与用户所在组织存在既定业务关系的各方，包括员工、承包商、供应商及其他业务合作伙伴；外部攻击者则是指与用户所在组织无此类关系的第三方。这两类攻击者采用相同的攻击机制——即利用伪装的进站消息（例如乍看无害但实际包含隐藏指令的电子邮件）进行攻击；然而，它们在攻击所依赖的上下文知识方面存在差异：内部攻击者掌握着用户工作环境的相关信息，包括组织架构、同事姓名、当前进行中的项目以及特定项目的上下文信息，这使得其攻击行为更具可信度；而外部攻击者仅能利用公开可获取的信息（如专业社交平台或公司官网）进行攻击。两类攻击者均无法直接访问代理程序的 API、修改内存存储区或入侵用户账户。

攻击者试图诱使智能体在其持久化存储器中存储恶意或误导性信息，从而影响其未来的行为。攻击者的攻击目标可分为以下四类：

- 1) **完整性破坏**：对手注入误导性信息，导致智能体在后续交互中产生错误输出或执行非预期操作。
- 2) **敏感信息泄露**：攻击者诱使代理方披露用户的隐私信息。
- 3) **隐蔽数据外泄**：攻击者植入指令，使代理程序秘密地将有关用户活动的信息传输给外部方。
- 4) **范围越界**：对手诱导代理执行超出其预期范围的操作。

4. GhostWriter架构

在本节中，我们将详细阐述 GhostWriter 的工作机制并逐步解析其形式化协议。GhostWriter 采用一个两阶段协议：**注入阶段**（第4.1节），即攻击者向目标系统注入受污染的内存；以及**激活阶段**（第4.2节），即通过良性查询检索出该受污染的内存，从而改变智能体的行为。请注意，攻击者仅控制注入的内容；他们对后续用户提示或智能体的检索机制不具有任何影响。图3展示了这两个阶段的总体流程：攻击者发送一封邮件，声称其同事的电子邮件地址已变更。该信息由智能体存储，随后当用户请求向其同事发送邮件时，智能体便会将邮件重定向至恶意地址。

4.1. 内存注入

GhostWriter 的第一阶段在协议1中正式描述；该协议始于对手 A 选择一个目标（g）（第...行）。




Figure 3: Two-phase memory injection attack: Phase 1 poisons the memory, which is retrieved and acted upon in Phase 2 during a legitimate user request for an email address.

Protocol 1 Memory Injection (Phase 1)

Require: Adversary ($\mathcal{A}$), Agent ($\mathcal{M}$), User ($\mathcal{U}$), ingestion function (f), memory store ($\mathcal{S}$).

Adversary $\mathcal{A}$

- 1: Select goal $g \in \{\text{Integrity_corruption, Sensitive_info_leakage, Covert_exfiltration, Scope_escalation}\}$;
- 2: $t \leftarrow \text{SelectTarget}(g)$; // Target prompt is chosen.
- 3: $p \leftarrow \text{Construct}(g, t)$; // Attack payload p is created.
- 4: $x \leftarrow \text{Embed}(p)$; // p is inserted into a email.
- 5: Deliver x to $\mathcal{U}$'s inbox;

Agent $\mathcal{M}$

- 6: $m \leftarrow f(x)$; store $m \in \mathcal{S}$ // Email is converted into a memory.
- Injection succeeds** $\iff p$ preserved within $m \in \mathcal{S}$

1). $\mathcal{A}$ then selects a target prompt t (e.g., “Send an email to the client”), that aligns with g (Line 2). Both g and t are then used to construct a malicious payload p , which is then embedded within a legitimate-looking message x , such as a team meeting email (Lines 3-4). $\mathcal{A}$ then sends x to the inbox of their targeted user ($\mathcal{U}$), at which point agent $\mathcal{M}$ will automatically pass the message to its ingestion function f (Line 5-6). The function returns a created memory m , and adds it to its memory store $\mathcal{S}$ (Line 6). If p is preserved within m , then Phase 1 is considered successful.

To increase the probability that the malicious memory is retrieved by the user’s agent, an adversary can optimize their payload to better match the targeted prompt. This optimization can take place in black-, gray-, or white-box settings. Similar to how attacks on machine learning models are studied. While we consider all of the settings to be relevant, as outlined by our security assumptions we consider only a black-box approach.

$$\text{Score}(p, T, E) = \frac{1}{|T|} \sum_{q_i \in T} \cos(E(p), E(q_i)) \quad (1)$$

Protocol 2 describes our approach, in which an adversary can use a public email corpus, such as the Enron Email Corpus [28], clustering it into k topics which are stored in Q , examples of these topics include client communications,

Protocol 2 Black-Box Retrieval Optimization

Require: Payload (p), email corpus ($\mathcal{C}$), embedding model (E), number of topics (k), target prompt (t), similarity threshold (τ).

Ensure: Optimized payload p^*

- 1: $Q \leftarrow \text{Cluster}(\mathcal{C}, k, E)$; // Cluster $\mathcal{C}$ into k topics.
- 2: $T \leftarrow \text{classify}(t, Q, E)$; // Find topic cluster for t .
- 3: $s \leftarrow \text{Score}(p, T, E)$; // Score p 's similarity to T ;
- 4: **while** $s < \tau$ **do**
- 5: $p' \leftarrow \text{Rewrite}(p)$;
- 6: $s' \leftarrow \text{Score}(p', T, E)$;
- 7: **if** $s' > s$ **then**
- 8: $p \leftarrow p'$, $s \leftarrow s'$; // Update p when p' scores better.
- 9: **end if**
- 10: **end while**
- 11: **return** $p^* = p$; // Return the final result.

deadline extensions, workflow policies, etc (Line 1). The adversary can then identify the cluster (T) most closely aligned to their target prompt t , allowing them to score (s) the similarity of their attack email to T (Lines 2-3). The score is calculated using Equation 1, which takes an attack payload p , a topic cluster T , and an embedding algorithm E as parameters. For E there are several options to choose from for the attacker [29], [30], which may or may not align with the user’s agent’s embedding model.

The computed similarity score s is then compared to a target threshold τ . If the score is below τ , then p will be rewritten and saved as a temporary attack email p' , which will then be scored and saved as s' (Lines 4-6). If the resulting score s' is higher than the previous score s , then p is overwritten with p' (Lines 7-8). This process is repeated until the similarity score exceeds τ , and the final optimized attack email is returned (Line 11). This iterative process ensures that the attack is potent.

4.2. Attack Activation

Protocol 3 outlines *GhostWriter*'s second phase. Once a poisoned memory m is added to the memory store $\mathcal{S}$, it lies dormant until it is later retrieved. The user ($\mathcal{U}$) submits

第一阶段：记忆注入

第二阶段：攻击激活

第1阶段 — 合法内存状态；第2阶段 — 注入阶段；第3阶段 — 毒化内存状态




图 3：双阶段内存注入攻击：第一阶段对内存进行污染；在第二阶段，当合法用户发起电子邮件地址查询请求时，对受污染的内存进行检索并执行相应操作。

试验方案1记忆注入（第一阶段）

需求：对手 (A)、智能体 (M)、用户 (U)、数据摄入功能 (f)内存存储器 (S)。

对手 A

- 1: 选择目标 $g \in \{\text{完整性破坏、敏感信息泄露、隐蔽数据外泄、权限提升}\}$;
- 2: $t \leftarrow \text{SelectTarget}(g)$; // 已选择目标提示。
- 3: $p \leftarrow \text{Construct}(g, t)$; // 攻击载荷 p 创建完成。
- 4: $x \leftarrow \text{Embed}(p)$; // p 被插入到一封电子邮件中。
- 5: 将 x 发送至 U 的收件箱;

代理 M

- 6: $m \leftarrow f(x)$; 将 $m \in S$ 存入内存 // 电子邮件被转换为内存。**注入操作成功** $t \cdot p$ 在 m 中保持不变

1) A 随后选择一个目标提示词 t (例如“向客户端发送一封电子邮件”), 该提示词与 g (第 2 行) 相匹配。接着, 利用 g 和 t 构建一个恶意负载 p , 然后将该恶意负载嵌入到一个看似合法的消息 x 中 (例如团队会议邮件, 见第 3-4 行)。 A 随后将 x 发送至目标用户 (U) 的收件箱; 此时, 代理 M 将自动将该消息传递给其数据摄入函数 f (第 5-6 行)。该函数返回一个新建的内存条目 m , 并将其添加到其内存存储库 S 中 (第 6 行)。若 p 被保留在 m 中, 则认为第一阶段执行成功。

为提高用户代理检索到恶意内存地址的概率, 攻击者可对攻击载荷进行优化, 使其更契合目标提示词; 这种优化可在黑盒、灰盒或白盒场景下进行。这与对机器学习模型进行攻击研究的方式类似。尽管我们认为所有上述场景均具有相关性, 但根据我们的安全假设, 我们仅考虑黑盒攻击方式。

$$\text{Score}(p, T, E) = \frac{1}{|T|} \sum_{q_i \in T} \cos(E(p), E(q_i)) \quad (1)$$

协议 2 描述了我们的攻击方法: 攻击者可利用公共电子邮件语料库 (例如 Enron Email Corpus[28]) 将其聚类为 k 个主题, 这些主题存储在 Q 中; 这些主题的示例包括

协议 2 黑盒检索优化

输入参数：有效载荷 (p)、电子邮件语料库 (C)、嵌入模型 (E)、主题数量 (k)、目标提示词 (t)、相似度阈值 (ϵ)。

确保：优化有效载荷 p^*

- 1: $Q \leftarrow \text{Cluster}(C, k, E)$; // 将聚类 C 分为 k 个主题。
- 2: $T \leftarrow \text{classify}(t, Q, E)$; // 为 t 寻找主题簇。
- 3: $s \leftarrow \text{Score}(p, T, E)$; // 计算 p 与 T 之间的相似度;
- 4: 当 $s < \epsilon$ 时, 执行
- 5: $p' \leftarrow \text{重写}(p)$;
- 6: $s' \leftarrow \text{分数}(p', T, E)$;
- 7: 若 $s' > s$ 则
- 8: $p \leftarrow p', s \leftarrow s'$; // 当 p' 的得分更高时, 更新 p 。
- 9: 结束 if
- 10: 结束 (while)
- 11: 返回 $p^* = p$; // 返回最终结果。

客户沟通、截止日期延期、工作流程政策等 (第 1 行)。随后, 攻击者可识别出与其目标提示 t 最匹配的聚类 T , 从而对其攻击邮件与 T 之间的相似度进行评分 s (第 2-3 行)。该评分通过公式 1 计算得出, 该公式将攻击载荷 p 、主题聚类 T 以及嵌入算法 E 作为参数。对于 E , 攻击者可从多种选项中进行选择[29]。[30]该结果可能与用户代理的嵌入模型相吻合, 也可能不相吻合。

计算得到的相似度得分 s 随后与目标阈值 ϵ 进行比较。若该得分低于 ϵ , 则 p 将被重写并保存为临时攻击邮件 p' , 随后对该邮件进行评分并保存为 s' (第 4-6 行)。若所得得分 s' 高于先前得分 s , 则 p 将被覆盖为 p' (第 7-8 行)。此过程将持续重复, 直至相似度得分超过 ϵ , 随后返回最终优化后的攻击邮件 (第 11 行)。该迭代过程可确保攻击邮件具有强效性。

4.2. 攻击激活

协议 3 阐述了 *GhostWriter* 的第二阶段。一旦受污染的内存 m 被添加到内存存储集 S 中, 该内存将处于休眠状态, 直至后续被检索。用户 (U) 提交

Protocol 3 Attack Activation (Phase 2)

Require: Agent ($\mathcal{M}$), User ($\mathcal{U}$), memory store ($\mathcal{S}$), retrieval window (k), poisoned memory (m).

```

    User  $\mathcal{U}$ 
1: Submits a benign prompt  $q$  to  $\mathcal{M}$ .
    Agent  $\mathcal{M}$ 
2:  $\mathcal{R} \leftarrow \text{Retrieve}(q, k, \mathcal{S})$ ; //  $k$  entries are retrieved from  $\mathcal{S}$ .
3:  $C \leftarrow \{q\} \cup \mathcal{R}$ ; // Retrieved memories  $R$  is added to context.
4:  $\mathcal{M}$  generates response  $r$  based on context  $C$ ;
5: if  $m \in \mathcal{R}$  then
6:    $r$  deviates from  $\mathcal{U}$ 's intent;
7: else
8:    $r$  proceeds normally; //  $\mathcal{U}$ 's intent is preserved.
9: end if
```

a prompt (q) to its agent ($\mathcal{M}$), which then passes q to its memory retrieval function (Line 1). The function returns a set of memories, $\mathcal{R}$, composed of the top k results that most closely align with q (Line 2). $\mathcal{R}$ is then added to the agent's current context (C), which is utilized to generate a response r (Lines 3-4). If m is present in $\mathcal{R}$, the agent's behavior will be steered toward the adversary's goal (Lines 5-6). Otherwise, the agent proceeds normally (Lines 7-9).

Once the user submits a prompt that closely mirrors the targeted prompt, such as 'send an email', 'schedule a meeting', etc., the agent will be highly likely to retrieve the poisoned memory from its memory store. The agent, believing the memory to be trusted, follows the malicious directive planted within, which may have it add a malicious address to an email or leak sensitive information. An agent's memory architecture affects its susceptibility to retrieving poisoned memories, as it may apply summarization steps, delete older memories in favor of newer ones, or combine details from a new memory into an old one. The impact of these architectural differences is analyzed in Section 7.

5. AM-Sentry Architecture

Proposed memory systems for agents largely ignore the adversarial threat vector introduced by storing inputs in long-term memory. These ingestion channels also give an adversary a path to inject false or malicious information into the agent's memory, which can then shape the summaries it produces, the recommendations it makes, and the actions it takes on the user's behalf. Adversaries can learn a lot about a company, such as employee names, communication protocol, allowing them to craft an attack payload indistinguishable from legitimate sources, as seen in targeted business email compromise (BEC) attacks [31].

In this context, spam filtering is insufficient because a targeted adversary can craft semantically valid messages that evade aggregate classifiers, and a single injection can persist across sessions. Defending against this threat requires policy directives defining which inputs may modify the long-term state. A secure memory governance system must be designed to satisfy the following goals:

- *Preserve utility*: maintain the assistant's ability to reason over accumulated context.

- *Signature-independent defenses*: defend against novel, previously unseen payloads.
- *Malicious memory characterization*: define what makes a memory unsafe.
- *Source trust modeling*: assess inputs by where they originate, not only what they contain.
- *Architecture-agnostic*: operate across assistant implementations and model families without assumptions about the underlying memory or model.
- *Two-stage coverage*: protection at memory saving and memory retrieval.

We introduce *AM-Sentry*, a modular, two-stage defense system for securing stateful agentic applications. *AM-Sentry* comprises two complementary mechanisms: (i) a policy component that governs which inputs are saved to memory, and (ii) a retrieval screening component that inspects retrieved memories before they influence the agent. Each component uses an LLM as a judge alongside non-LLM decision logic. In the remainder of this section, we describe each mechanism in detail and demonstrate how their combination satisfies the outlined design goals.

5.1. Policy Configurations

Current agentic memory systems store inputs without restriction, regardless of future utility or safety, allowing malicious inputs to enter long-term memory. A policy that determines what is stored closes this vulnerability, but it requires a structured memory representation so that each input can be judged uniformly. The policy imposes this structure and evaluates every candidate against a set of rules before admission to the memory store. Figure 4 illustrates this with an example of our policy blocking a malicious email from being saved into the memory store. We formalize *AM-Sentry* with three distinct policy configurations (S_1 , S_2 , and S_3) of varying strictness. Moving from S_1 to S_3 , each configuration builds on the security properties of the previous through tighter rules, a richer memory structure, and more rigorous security assessment.




Figure 4: Incoming content is evaluated before being admitted to the memory store. Legitimate updates are accepted, while malicious content is blocked.

协议3攻击激活（第二阶段）

需要：代理（ M ）、用户（ U ）、内存存储器（ S ）、检索窗口（ k ）、中毒内存（ m ）。

用户 U

1: 将良性提示 q 提交给 M 。

代理 M

2: $R \leftarrow \text{Retrieve}(q, k, S)$; // 从 S 中检索 k 条目。

3: $C \leftarrow \{q\} \cup R$; // 检索到的内存 R 被添加到上下文之中。

4: M 根据上下文 C 生成响应 r ;

5: 若 $m \in \mathbb{R}$, 则

6: r 与 U 的初衷相悖;

7: 否则

8: r 正常运行; // U 的意图得以保留。

9: 结束 if

一个发送给其智能体（ M ）的提示词（ q ），随后该智能体将 q 传递给其记忆检索函数（第1行）。该函数返回一组记忆 R ，该集合由与 q 最匹配的前 k 个结果组成（第2行）。接着将 R 添加到智能体的当前上下文（ C ）中，该上下文用于生成响应 r （第3-4行）。若 m 存在于 R 中，则智能体的行为将被引导至对抗方的目标方向（第5-6行）；否则，智能体按常规流程执行（第7-9行）。

当用户提交的提示语与目标提示语高度相似（例如“发送一封电子邮件”、“安排一次会议”等）时，智能体极有可能从其记忆库中检索到被污染的记忆。由于智能体认为该记忆可信，便会执行其中植入的恶意指令——例如在电子邮件中添加恶意地址或泄露敏感信息。智能体的记忆架构会影响其检索到被污染记忆的易受攻击程度；例如，智能体可能会应用摘要化处理步骤、删除旧记忆以保留新记忆，或将新记忆中的细节合并到旧记忆中。这些架构差异所带来的影响将在第7节中进行分析。

5. AM-Sentry 架构

现有的智能体内存系统在很大程度上忽略了因将输入数据存储在长期记忆中而引入的对抗性威胁向量。这些数据摄入通道还为攻击者提供了向智能体内存中注入虚假或恶意信息的途径；此类信息随后可影响智能体生成的摘要、所做出的推荐以及代表用户执行的操作。攻击者能够深入了解企业相关信息（例如员工姓名、通信协议），从而使他们能够构建与合法来源无异的攻击载荷——这在针对性企业电子邮件入侵（BEC）攻击中可见一斑[31]。

在此背景下，传统的垃圾邮件过滤机制已显不足——因为恶意攻击者能够生成语义上有效的消息以绕过聚合分类器，且单次注入攻击可在多个会话间持续存在。防范此类威胁需要制定相应的策略指令，明确哪些输入内容可能影响系统的长期状态。因此，必须设计一套安全的内存治理系统，以实现以下目标：

- **保持实用性强：**维持助手基于累积上下文进行推理的能力。

- **独立于签名的防御机制：**用于防御新型、此前未见过的恶意载荷。
- **恶意内存特征分析：**定义导致内存不安全的原因。
- **源信任建模：**根据输入数据的来源进行评估，而非仅依据其包含的内容。
- **与架构无关：**可跨不同助手实现及模型系列进行操作，无需对底层内存或模型作任何假设。
- **两级覆盖：**在内存保存与内存检索阶段均提供保护。

我们推出 *AM-Sentry*——一款用于保护具有状态感知特性的智能体应用程序的模块化、两阶段防御系统。*AM-Sentry* 包含两种互补的机制：（i）一个用于管控哪些输入数据应保存至内存的策略组件；以及（ii）一个在检索到的内存数据影响智能体之前对其进行审查的检索筛选组件。每个组件均采用大型语言模型（LLM）作为判定依据，并结合非 LLM 决策逻辑共同运作。在本节的剩余部分中，我们将详细阐述每种机制，并演示这两种机制的结合如何满足上述设计目标。

5.1. 策略配置

当前的代理式内存系统对输入数据进行无限制存储，无论其未来是否具有实用价值或安全性，这使得恶意输入数据能够进入长期内存。通过制定一项决定存储内容的策略可消除这一漏洞，但该策略需要一种结构化的内存表示方式，以便对每个输入数据进行统一评估。该策略为内存存储系统施加了此类结构，并在允许将输入数据存入内存存储之前，根据一组规则对每个候选数据进行评估。图4通过一个示例展示了这一机制——即我们的策略阻止了一封恶意电子邮件被保存到内存存储中。我们针对 *AM-Sentry* 定义了三种严格程度不同的策略配置（ S_1, S_2 及 S_3 ）。从 S_1 到 S_3 的演进过程中，每种配置均在前一配置的安全特性基础上进行扩展，具体表现为采用更严格的规则、构建更丰富的内存结构以及实施更严谨的安全性评估。




图 4：入站内容在被存入内存存储器之前会经过评估；合法的更新内容将被接受，而恶意内容则会被拦截。

Protocol 4 S_2 Policy Configuration

Require: Candidate memory (m), LLM judge ($\mathcal{J}$).**Ensure:** Admit or reject m **LLM Judge $\mathcal{J}$**

- 1: Scores m on four fields (each 0.0–1.0);
- 2: $o \leftarrow \text{Origin}(m)$; // source identifiability
- 3: $r \leftarrow \text{Trust}(m)$; // reliability of information
- 4: $t \leftarrow \text{Type}(m)$; // fact | workflow | contact | instruction
- 5: $u \leftarrow \text{Utility}(m)$; // future task relevance

Non-LLM Decision Logic

- 6: $U \leftarrow 0.6 \cdot u + 0.4 \cdot r$ // usefulness equation
 - 7: $V \leftarrow t \times (1 - o)$; // vulnerability equation
 - 8: **if** $U > 0.4$ **and** $V < 0.6$ **then**
 - 9: **return** ADMIT;
 - 10: **else**
 - 11: **return** REJECT;
 - 12: **end if**
-

Each policy configuration uses an LLM judge to evaluate a memory candidate across varying metrics and formats. The first policy configuration, S_1 , consists of a single LLM judge that determines whether a memory candidate should be saved based on the potential future use of the information in the proposed memory. S_1 has the LLM judge make a simple yes-or-no decision.

Our second configuration, S_2 , introduces a more explicit structure and is formalized in Protocol 4. An LLM judge evaluates memory candidates across four fields, which are described in Table 1. Each field is assigned a numerical score between 0 and 1 (Lines 1-5). These fields are then passed to the non-LLM decision logic and used to make a final determination on the memory. The utility (u) and trust (r) fields are used to determine a usefulness score (U) (Line 6). The type (t) is used in an equation with the origin field (o) to determine a vulnerability score (V) (Line 7). The memory is admitted only if U is sufficiently high and V is sufficiently low (Lines 8-12).

S_3 is the strictest configuration; every memory candidate is evaluated on a seven-point checklist, listed in Table 2. There are three critical items (*vulnerability*, *sensitivity*, and *source trust*), with the *vulnerability* triggering a hard rejection if its score falls below the failing threshold. S_3 is formalized in Protocol 5, which begins with an LLM judge that evaluates each checklist item and assigns a score between 0 and 1 (Lines 1-2). The filled-in checklist is then passed to non-LLM decision logic, which first checks

TABLE 1: S_2 memory fields. An LLM judge evaluates all four fields and assigns score to each.

Field	Scoring Guidance
Origin	User statement (1.0), Internal email/calendar (0.8), Client email/calendar (0.5), Untrusted email/calendar (0.2), or Unknown type (0.0).
Trust	High (1.0), Medium (0.5), or Low (0).
Type	Instruction (1.0), Workflow (0.7), Preference (0.5), Contact information (0.1), or Fact (0.0).
Utility	Essential (1.0), Moderately useful (0.5), Not useful (0.0).

Protocol 5 S_3 Policy Configuration

Require: Candidate memory (m), source metadata (SM), LLM judges ($\mathcal{J}_1, \mathcal{J}_2$).**Ensure:** Admit or reject m **First LLM Judge $\mathcal{J}_1$**

- 1: Scores m on seven-point checklist (each 0.0–1.0);
- 2: $s_1 \leftarrow \text{Vulnerable}$; $s_2 \leftarrow \text{Sensitivity}$; $s_3 \leftarrow \text{Source}$;
 $s_4 \leftarrow \text{Contradict}$; $s_5 \leftarrow \text{Instruction}$; $s_6 \leftarrow \text{Verify}$;
 $s_7 \leftarrow \text{Justify}$;

Non-LLM Decision Logic

- 3: **if** $s_1 < 0.6$ **then**
- 4: **return** REJECT; // hard reject on vulnerability
- 5: **end if**
- 6: $C \leftarrow (s_2 + s_3) / 2$; // Critical item equation
- 7: $N \leftarrow (s_4 + s_5 + s_6 + s_7) / 4$; // Non-critical item equation
- 8: $F \leftarrow 0.6 \cdot C + 0.4 \cdot N$; // Final score equation
- 9: **if** $F < 0.4$ **then**
- 10: **return** REJECT;
- 11: **end if**

Second LLM Judge $\mathcal{J}_2$

- 12: $\mathcal{J}_2$ receives SM and s_1 - s_7 .
 - 13: **return** $\mathcal{J}_2$'s admit/reject decision. // Looks for contradictions
-

whether the vulnerability score is too low and rejects it if so (Lines 3-5). Otherwise, the other two critical items and the non-critical items are combined to determine the final score (F) for admission (Lines 6-8). If F falls below the acceptance threshold, the memory is rejected. Otherwise, the scores and memory metadata are passed to a second LLM judge, which looks for signs that the first LLM judge was compromised through prompt injection (Lines 9-11). If the second judge reasonably suspects the first judge was compromised, then the memory is dropped (Lines 12-13).

Our policy configurations satisfy our design goals. They *preserve utility* by filtering out content deemed unnecessary, or that clearly signals malicious intent, thereby allowing legitimate useful content to be admitted. Our policy configurations do not rely on pattern matching known attacks, making them *signature-independent defense* that makes per-memory decisions based on structure, provenance, and risk. S_3 's vulnerability assessment directly tests whether content could redirect communications, enable unauthorized actions, or leak information, providing *malicious memory charac-*

TABLE 2: S_3 checklist. C = critical, NC = non-critical.

Item	Level	Scoring Guidance
Vulnerability	C	No harm potential (1.0), Could change agent behavior (0.0).
Sensitivity	C	No sensitive data discussed (1.0), Overrides sensitive data. (0.0).
Source trust	C	User (1.0), Internal contact (0.8), Client (0.3), External untrusted (0.0).
Contradiction	NC	No conflicts (1.0), Contradicts multiple other memories (0.0).
Instruction	NC	Informative (1.0), Directives (0.0).
Verifiability	NC	Verifiable (1.0), Unverifiable (0.0).
Justification	NC	Clear future use (1.0), Seemingly useless information (0.0).

协议 4 S_2 策略配置

需要：候选人记忆 (m)、LLM 评估器 (J)。

确认：接收或拒绝 m

LLM法官 J

- 1: 四个领域上的得分 m (每个领域评分范围为 0.0–1.0) ;
- 2: $o \leftarrow$ 原点 (m); // 源可识别性
- 3: $r \leftarrow \text{Trust}(m)$; // 信息可靠性
- 4: $t \leftarrow \text{Type}(m)$; // 事实 | 工作流 | 联系人 | 指令
- 5: $u \leftarrow$ 效用函数(m); // 未来任务相关性

非大语言模型决策逻辑

- 6: $U \leftarrow 0.6 \cdot u + 0.4 \cdot r$ // 有效性方程
- 7: $V \leftarrow t \times (1 - o)$; // 漏洞方程
- 8: 若 $U > 0$ 且 $V < 0.6$, 则
- 9: 返回录取;
- 10: 否则
- 11: 返回拒绝;
- 12: 结束 if

每项策略配置均使用一个大型语言模型 (LLM) 评估器，根据不同的衡量指标和格式对记忆候选项进行评估。第一项策略配置 S_1 包含一个单一的 LLM 评估器，该评估器根据所提议记忆中信息在未来的潜在用途来决定是否应保存该记忆候选项。 S_1 由该 LLM 评估器做出简单的“是”或“否”型判定。

我们的第二种配置 S_2 引入了更为明确的结构，并在协议 4 中进行了形式化描述。LLM 评估器会对内存候选项进行跨四个字段的评估，这些字段在表 1 中有详细说明。每个字段被赋予一个介于 0 到 1 之间的数值评分 (第 1-5 行)。随后，这些字段被传递给非 LLM 决策逻辑，用于对内存进行最终判定。效用 (u) 和信任度 (r) 字段用于计算一个实用性评分 (U) (第 6 行)。类型 (t) 与源字段 (o) 共同参与一个方程，用于计算漏洞评分 (V) (第 7 行)。仅当 U 足够高且 V 足够低时，该内存才被允许使用 (第 8-12 行)。

S_3 是最严格的配置方案；每个内存候选项均需通过一份包含七个项目的检查清单进行评估，该清单列于表 2 中。该方案包含三项关键项 (漏洞、敏感性 及 源代码可信度)，若漏洞的得分低于失败阈值，则触发硬拒绝机制。 S_3 已在协议 5 中正式定义；该协议由一名 LLM 评审员负责对每个检查清单项进行评估并赋予 0 到 1 之间的分数 (第 1-2 行)，随后将填写完整的检查清单传递给非 LLM 决策逻辑模块，该模块首先进行检查

表 1: S_2 存储字段。LLM 评审员对所有四个字段进行评估，并为每个字段赋分。

场地	评分指南
Origin	用户声明 (1.0)、内部电子邮件/日历 (0.8)、客户端电子邮件/日历 (0.5)、不可信电子邮件/日历 (0.2) 或未知类型 (0.0)。
Trust 类	高 (1.0)、中 (0.5) 或低 (0)。
型公用	指令 (1.0)、工作流 (0.7)、偏好设置 (0.5)、联系方式 (0.1) 或事实 (0.0)。
事业服	非常有用 (1.0)、中等有用 (0.5)、无用 (0.0)。
务	

协议 5 S_3 策略配置

要求：候选人记忆 (m)，源判别器 ($\mathcal{J}$ 元数据 (SM)、大语言模型 (LLM) $1, \mathcal{J}_2$)。

确认：接收或拒绝 m

首任大型语言模型法官 $\mathcal{J}_1$

- 1: Scores m on seven-point checklist (each 0.0–1.0);
- 2: $s_1 \leftarrow$ 漏洞; $s_2 \leftarrow$ 敏感性; $s_3 \leftarrow$ 源; $s_4 \leftarrow$ 矛盾; $s_5 \leftarrow$ 指令; $s_6 \leftarrow$ 验证; $s_7 \leftarrow$ 证明;

非大语言模型决策逻辑

- 3: 若 $s_1 < 0$ 且 6, 则
- 4: 返回 reject; // 对漏洞进行硬拒绝
- 5: 结束 if
- 6: $C \leftarrow (s_2 + s_3) / 2$; // 关键项方程
- 7: $N \leftarrow (s_4 + s_5 + s_6 + s_7) / 4$; // 非关键项方程
- 8: $F \leftarrow 0.6 \cdot C + 0.4 \cdot N$; // 最终得分方程
- 9: 若 $F < 0.4$ 则
- 10: 返回拒绝;
- 11: 结束 if

第二名大型语言模型法官 $\mathcal{J}_2$

- 12: $\mathcal{J}_2$ 接收 SM 与 s_1-s_7 。
- 13: 返回 $\mathcal{J}_2$ 的接受/拒绝决策。// 检测矛盾之处

若漏洞评分过低，则予以拒绝 (第 3–5 行)；否则，将另外两项关键项与非关键项合并，以确定用于准入的最终评分 (F) (第 6–8 行)。若 F 低于接受阈值，则该内存数据被拒绝；否则，将相关评分及内存元数据传递给第二位大型语言模型 (LLM) 评审员，由其检测是否存在第一位 LLM 评审员因提示词注入而遭到攻破的迹象 (第 9–11 行)。若第二位评审员合理怀疑第一位评审员遭到攻破，则该内存数据将被丢弃 (第 12–13 行)。

我们的策略配置符合我们的设计目标。这些配置通过过滤掉被判定为不必要的内容或明显带有恶意意图的内容来保障系统效用，从而允许合法且有用的内容通过。我们的策略配置不依赖于对已知攻击模式的匹配，因此属于一种与特征无关的防御机制，其基于内容结构、来源及风险做出永久性内存决策。 S_3 的漏洞评估直接检测内容是否可能重定向通信、引发未经授权的操作或泄露信息；该评估提供了恶意内存特征表 2: S_3 检查清单。C = 严重，NC = 非严重。

项目	级别	评分指南
漏洞	C	无危害性 (1.0)，可能改变代理人行为 (0.0)。
灵敏度	C	无敏感数据涉及 (1.0)，覆盖敏感数据 (0.0)。
源信任	C	用户 (1.0)、内部联系人 (0.8)、客户端 (0.3)、外部不可信 (0.0)。
矛盾	NC	无冲突 (1.0)，与多个其他记忆相矛盾 (0.0)。
指令	NC	信息类 (1.0)，指令类 (0.0)。
可验证性	NC	可验证 (1.0)，不可验证 (0.0)。
理由	NC	明确的未来用途 (1.0)，看似无用的信息 (0.0)。




Figure 5: The retrieval screen blocks poisoned memories before they can enter an agent’s context.

terization. S_2 and S_3 make use of *source trust modeling*. Our policy is *architecture-agnostic* as it intercepts memory candidates before they enter a given agent’s memory store, and does not require changes to the underlying agent.

5.2. Memory Retrieval Screen

When poisoned content is saved in the memory store, the defense must intercept it at retrieval time. We accomplish this with a retrieval screen that evaluates each retrieved memory against the current context and blocks unsafe memories before they reach the agent. Figure 5 gives an example of the retrieval screen rejecting a suspicious memory containing an instruction to redirect communications. Along with the memory-saving policy, the retrieval screen provides in-depth defense against memory-injection attacks, ensuring that any memory-bypassing admission is re-evaluated before it influences the agent. When paired with either S_2 or S_3 , the screen can leverage the structured metadata they produce, enabling it to operate on explicit trust, risk, and type annotations (e.g., `[trust=external-untrusted, risk=medium, type=instruction]`) rather than inferring them from memory itself.

The screen first evaluates each retrieved memory against the current query. The screen considers four criteria, summarized in Table 3: *relevance*, *instruction suppression*, *contradiction detection*, and *source trust*. An LLM judge is asked to evaluate the first three criteria to determine the utility and safety of a retrieved memory. The LLM judges and considers those criteria in the context of source trust; for example, a memory containing a directive will be retained if the source is trusted, but is otherwise discarded. After iterating through all memories, the LLM judge returns the surviving memories to the agent’s context.

The retrieval screen aims to filter malicious memories while preserving utility. It *preserves utility* by minimizing unnecessary exclusions and admitting legitimate content into the agent’s context. It provides *malicious memory characterization* by evaluating each candidate against known attack properties, specifically embedded instructions and unverifiable external claims. It *models source trust* by conditioning

TABLE 3: Retrieval screen exclusion criteria. The screen drops a candidate if any criterion fires.

Criterion	Exclusion rule
Relevance	Memory not directly relevant to the current query.
Instruction suppression	Memory contains directives or behavioral rules; memories are data, not commands.
Trust filter	Memory from an untrusted external source making unverifiable claims.
Contradiction detection	Contradicts existing trusted information.

on memory origin during evaluation, exploiting S_2 and S_3 metadata annotations when available. Finally, it completes *two-stage coverage*, serving as the memory-retrieval counterpart to the memory-saving policy and intercepting attacks that survive admission.

6. Experimental Methodology

6.1. Agent Selection

Given the absence of long-term memory agents that are purpose-built for a personal assistant use case, we selected five state-of-the-art agent architectures and adapted them to a common evaluation setting: A-Mem [22], Mem0 [32], ExpeL [4], Letta (formerly MemGPT) [3], [33], and MemoryOS [5]. These agents represent a variety of memory representation, organization, and retrieval techniques, allowing us to compare against a variety of design choices. The majority of selected agents use fact-based memory, as they seemingly suit the use case best. We also included ExpeL, an agent with trajectory-based memory that stores successful action plans rather than facts, introducing distinct vulnerability characteristics for us to assess.

We built a wrapper for each agent that exposes a common interface. We only made architectural adjustments to allow us to evaluate each agent uniformly while preserving the mechanisms central to each agent’s design. The wrapper includes a system prompt that informs the agent of their




图 5：检索筛选机制可在受污染的记忆进入智能体上下文之前将其拦截。

特征提取。 S_2 与 S_3 采用**源信任建模**。我们的策略具有**架构无关性**，因为它在内存候选项进入指定智能体的内存存储器之前对其进行拦截，且无需对底层智能体进行任何修改。

5.2. 记忆提取界面

当中毒内容被保存到内存存储器中时，防御机制必须在检索时拦截该内容。我们通过一个检索筛选器来实现这一功能——该筛选器会将每个检索到的内存条目与当前上下文进行比对，并在这些内存条目传递至智能体之前将其拦截。图5展示了该检索筛选器对一条包含“重定向通信”指令的可疑内存条目的拦截示例。除内存保存策略外，该检索筛选器还提供了针对内存注入攻击的深度防御机制，确保任何绕过内存保护的访问行为在影响智能体之前均经过重新评估。当与 S_2 或 S_3 结合使用时，该筛选器可利用它们生成的结构化元数据，从而基于显式信任、风险及类型标注（例如：`[trust=external-untrusted, risk=medium, type=instruction]`）进行操作，而非从内存本身推断这些标注。

该筛选机制首先将每条检索到的记忆与当前查询进行比对。该机制考虑四个评估标准，如表3所列：**相关性**、**指令抑制**、**矛盾检测**以及**来源可信度**。由LLM评估员对前三个标准进行评估，以确定所检索记忆的实用性和安全性；该LLM评估员在评估时会结合来源可信度这一背景因素——例如，若来源可信，则保留包含指令的记忆，否则将其丢弃。在遍历所有记忆后，LLM评估员将筛选出的存活记忆返回至智能体的上下文环境中。

检索筛选界面旨在过滤恶意记忆的同时保留其效用。该界面通过最大限度减少不必要的内容排除并将合法内容纳入智能体的上下文之中，从而**保持效用**。它通过将每个候选内容与已知的攻击特征（特别是嵌入式指令和无法验证的外部声明）进行比对，来提供**恶意记忆特征分析**。该界面通过调整表3：检索筛选界面排除标准来**建模源信**

任度；若任何排除标准触发，则该候选内容将被剔除。

准则	排除规则
相关性指令抑制信任过滤器	内存数据与当前查询无关。
矛盾检测	内存包含指令或行为规则；内存存储的是数据，而非命令。来自不可信外部源的内存可能包含无法验证的声明。
	与现有可信信息相矛盾。

在评估过程中利用内存来源信息，并在条件允许时利用 S_2 与 S_3 元数据注解。最后，该机制实现**两阶段覆盖**，作为内存节省策略的对应机制，用于拦截那些在准入后仍持续存在的攻击行为。

6. 实验方法学

6.1. 代理选择

鉴于目前尚缺乏专门为个人助理应用场景设计的长时记忆智能体，我们选取了五种最先进的智能体架构，并将其适配至一个通用的评估环境：**A-Mem**[22]、**Mem0**[32]、**ExpeL**[4]、**Letta**（原名 **MemGPT**）[3]。[33]此外还有**MemoryOS**[5]。这些智能体代表了多种记忆表征、组织及检索技术，使我们能够对不同的设计方案进行比较。所选的大多数智能体均采用基于事实的记忆机制，因为它们似乎最适用于当前的用例；我们还纳入了**ExpeL**——一款采用基于轨迹的记忆机制的智能体，其存储的是成功的行动计划而非事实，这为我们的评估引入了独特的脆弱性特征。

我们为每个智能体构建了一个封装层，该封装层提供了统一的接口。我们仅对系统架构进行了调整，以便在保留各智能体设计核心机制的前提下，对所有智能体进行统一评估。该封装层包含一个系统提示语，用于告知智能体

assignment and capabilities, a common tool interface accessible to all agents, and standardization of input feeding and output gathering. Our simulated tool APIs allow agents to accept input from email or calendar events and sending email. The agents were instrumented to capture all their output and operations to support analysis and iterative refinement of the adaptations. *AM-Sentry* was also integrated into the wrapper, so that regardless of how an agent represents their memory, *AM-Sentry* could be applied to all.

These agents have various ways of saving user inputs as memories; some extract information and save only facts, others summarize the event and save only the summary, while others save the entire interaction, such as Letta. For most, we did not alter their design, though in the case of MemoryOS, a change was required. MemoryOS saves a conversation pair; the input is saved verbatim, and so is the agent’s response. The issue with that design is that many tasks result in a simple acknowledgment message from the agent, which creates problems with memory embeddings. To account for bare-bones responses, we changed MemoryOS to summarize its actions and tool inputs, better distinguishing memories from one another.

6.2. Memory Snapshot and Utility Test Suite

Each agent is loaded with a memory snapshot constructed from events representing emails, calendar events, and meetings spanning five days, and fed to the agents in chronological order. We list examples of these inputs in Appendix C. Some events were also intentionally designed to change or evolve past events, such as deadline extensions. Each agent builds their own memory store, saving information on contacts, deadlines, meetings, and emails. After processing the simulated work week, we saved a snapshot of the memory state for later reuse, allowing each agent to make use of its snapshot during the experiments.

To analyze the impact of our defenses, a baseline utility needed to be established for each agent, but no existing benchmark is well-suited to our use case. The closest are LoCoMo [34], HotpotQA [35], and Webshop [36], but all lack one or more aspects we sought to evaluate. None of these benchmarks focuses on evolving information, such as changing meeting times or contact information updates. Additionally, all information comes from a single source, and in the case of LoCoMo and HotpotQA, there are no tool-based evaluations, while Webshop is not fact-focused. A proper evaluation for these types of agents needs to test three things: multiple input sources, tool-based task evaluation, and events that evolve or change. To address this, we also constructed an evaluation based on the synthetic workweek. This allows us to determine the baseline utility and compare it with the utility scores for each *AM-Sentry* configuration. For our evaluation, we defined four categories: recall of static events, recall of evolved (changed) events, task execution using static information, and task execution using evolved information. Based on this, we can get a complete picture on agent utility and the impacts of the proposed defense mechanism.

7. Evaluation

7.1. Experiment Setup

We ran each agent with four LLM models: GPT-5.4-mini, DeepSeek-V4-Flash, Gemini-2.5-Flash, and Llama 3.1 8B. For brevity, throughout the remainder of this section, GPT-5.4-mini, DeepSeek-V4-Flash, Gemini-2.5-Flash, and Llama-3.1-8B are referred to as ChatGPT, DeepSeek, Gemini, and Llama, respectively, unless otherwise specified. We set the model temperature to 0.7, where applicable. All agents use their default configurations. We ran all experiments on a workstation with a 13th Gen Intel Core i9-13900K, 128 GB of RAM, and an NVIDIA RTX 4080 (16 GB VRAM). We host Llama locally via Ollama using the Q4_K_M quantization, and we access ChatGPT, DeepSeek, and Gemini through their respective APIs. We use each provider’s default decoding parameters across all models. All of our agents use the all-MiniLM-L6-v2 embedding model [29], while the adversary uses bge-small-en-v1.5 [30]. Each experiment was repeated five times to account for random variation unless otherwise noted.

7.2. GhostWriter Analysis

Experiment Design: We created 16 attack scenarios, varying the delivery method between email and calendar invite. For each test scenario, we used one malicious payload and tested the benign prompt once against agents with no active defense. Each scenario was tested five times, and we reported the aggregated results. We also tested two alternatives for each scenario: **directive payloads**, which have a more direct tone, and **descriptive payloads**, which describe false facts or make polite requests. Examples of each are in Appendix B; we mainly focus on the descriptive payload in this section. For each attack scenario, the agents’ memories are reset to a snapshot of a five-day workweek built by each agent, which is free of malicious memories.

We evaluate Phase 1 (P1) and deem the attack successful if the malicious payload is saved in memory. We also evaluate Phase 2 (P2) of the attack, using a benign trigger to determine the malicious memory’s *retrieval rate* and *activation rate* after retrieval.

Result Analysis: Our results are shown in Figure 6, where each bar cluster is organized by LLM model used by the agents, and each bar represents a memory agent architecture. Memory Injection (P1) indicates whether the adversarial payload was stored in memory. We observed 100% attack success across most agents, with an average of 98% across all agents, regardless of the LLM model used. The only agents with an attack success rate below 100% were Mem0 and ExpeL, which transform the raw content before saving it. As a result, attack payloads are not always preserved in memory for those two agents, resulting in lower attack success rates. We also see that attack retrieval is high across all agents when paired with ChatGPT and DeepSeek, ranging from 79-100% and averaging around 94% on these

其任务分配与能力范围；一个可供所有智能体访问的通用工具接口；以及对输入数据馈送与输出数据采集过程的标准化处理。我们设计的模拟工具 API 允许智能体接收来自电子邮件或日历事件的输入，并发送电子邮件。各智能体均经过配置，用于捕获其所有输出结果与操作日志，以支持对智能体适配方案进行分析与迭代优化。*AM-Sentry* 亦集成至该封装层中，因此无论智能体如何表示其记忆状态，*AM-Sentry* 均可应用于所有智能体。

这些智能体采用多种方式将用户输入保存为记忆：有些智能体会提取信息并仅保存事实，有些则会对事件进行摘要后仅保存摘要，而另一些（如 Letta）则会保存完整的交互过程。对于大多数智能体，我们并未对其设计进行修改；不过对于 MemoryOS，我们进行了相应的调整。MemoryOS 会保存一对对话记录——用户输入和智能体的响应均按原样保存。这种设计存在的问题在于：当执行某些任务时，智能体仅会返回一条简单的确认消息，这会对记忆嵌入表示法的处理造成困扰。为适配这种极简化的响应方式，我们对 MemoryOS 进行了修改，使其能够对自身操作及工具输入进行摘要处理，从而更清晰地区分不同的记忆条目。

6.2. 内存快照与实用程序测试套件

每个智能体均加载了一个由代表电子邮件、日历事件及会议的事件所构建的记忆快照（这些事件涵盖连续五天的数据），并按时间顺序将该快照输入至各智能体。我们已在附录C中列举了此类输入数据的示例。部分事件在设计时特意被设定为对过往事件进行修改或演变（例如截止日期延长）。每个智能体均构建了独立的记忆存储库，用于保存有关联系人、截止日期、会议及电子邮件的信息。在完成对模拟工作周的处理后，我们保存了当前记忆状态的快照以供后续复用，从而使各智能体在实验过程中能够调用其对应的快照。

为分析我们防御机制的影响，需为每个智能体建立一个基准效用值，但现有基准测试均不完全适用于我们的应用场景。最接近的基准测试包括 LoCoMo[34]、HotpotQA[35]和 Webshop[36]，但这些基准测试均缺乏我们想要评估的一个或多个方面。这些基准测试均未聚焦于动态变化的信息（例如会议时间变更或联系人信息更新）；此外，所有信息均源自单一来源；而在 LoCoMo 和 HotpotQA 的案例中，缺乏基于工具的任务评估；而 Webshop 则不以事实信息为中心。针对此类智能体的恰当评估需涵盖三个方面：多输入源、基于工具的任务评估以及动态变化或演进的事件。为解决这一问题，我们还构建了一个基于合成工作周的评估方案。这使我们能够确定基准效用值，并将其与每个 *AM-Sentry* 配置对应的效用得分进行比较。在本次评估中，我们定义了四个类别：静态事件的召回率、动态（变更）事件的召回率、基于静态信息的任务执行，以及基于动态信息的任务执行。基于此，我们可以全面了解代理效用以及所提出的防御机制所带来的影响。

7. 评估

7.1. 实验设置

我们分别使用四种大语言模型（LLM）对各智能体进行测试：GPT-5.4-mini、DeepSeek-V4-Flash、Gemini-2.5-Flash 以及 Llama 3.1 8B。为简洁起见，在本节后续内容中，除非另有说明，上述四种模型将分别简称为 ChatGPT、DeepSeek、Gemini 和 Llama。在适用情况下，我们将模型温度设置为 0.7；所有智能体均采用其默认配置。所有实验均在一台配备第13代 Intel Core i 9-13900K 处理器、128 GB 内存及 NVIDIA RTX 4080 显卡（16 GB 显存）的工作站上运行。我们通过 Ollama 在本地部署 Llama 模型（采用 Q4-K-M 量化方案），并通过 ChatGPT、DeepSeek 和 Gemini 的相应 API 进行调用。在所有模型中，我们均采用各提供商的默认解码参数。所有智能体均使用 all-MiniLM-L6-v2 嵌入模型 [29]，而对抗智能体则使用 bge-small-en-v1.5 [30]。除非另有说明，所有实验均重复五次以消除随机性影响。

7.2. GhostWriter 分析

实验设计：我们设计了 16 种攻击场景，其中攻击方式在电子邮件与日历邀请之间进行切换。对于每个测试场景，我们使用一个恶意载荷，并对未启用任何主动防御机制的智能体进行一次良性提示测试。每个场景均测试五次，随后汇总并报告结果。此外，我们为每个场景测试了两种替代方案：**指令型载荷**（语气更为直接）和**描述性载荷**（描述虚假事实或提出礼貌性请求）；各类载荷的示例详见附录B；本节主要聚焦于描述性载荷。对于每个攻击场景，智能体的内存均重置为由各智能体自构建的为期五天工作周的内存快照，该快照不含任何恶意内存痕迹。

我们对第一阶段（P1）进行评估；若恶意载荷成功保存于内存中，则判定此次攻击为成功。我们还对攻击的第二阶段（P2）进行评估，通过使用良性触发器来确定恶意内存的检索率以及激活率。

结果分析：我们的结果如图6所示；图中每个柱状簇根据智能体所使用的LLM模型进行分类，每根柱状条代表一种记忆智能体架构。“Memory Injection (P1)”表示对抗性载荷是否被存储在内存中。我们观察到，绝大多数智能体的攻击成功率均达到100%，所有智能体的平均攻击成功率为98%，且与所使用的LLM模型无关。唯一攻击成功率低于100%的智能体是Mem0和ExpeL——这两款智能体在保存内容前会对原始内容进行转换；因此，对于这两款智能体而言，攻击载荷并非总能保留在内存中，从而导致其攻击成功率较低。我们还发现，当与 ChatGPT 和 DeepSeek 配合使用时，所有智能体的攻击检索率均较高，其范围在79%–100%之间，平均约为94%。



Figure 6: *GhostWriter*'s success results, broken down into memory injection (P1) success, attack retrieval success, and attack activation success (P2).




Figure 7: P2 activation rate (%) by category and separated by model and agent.

models. Mem0 and ExpeL tend to have lower attack retrieval rates than the other agents, with rates as low as 57% and 65% for Gemini and Llama, respectively. This is because Gemini and Llama are less accurate when transforming the input into the memory, reducing the likelihood of matching that attack payload to the targeted benign prompt.

Attack activation is high across all models, averaging 60% across all agents and LLM models. ExpeL had the lowest activation rates, ranging from 28% to 63% across all LLMs, making it the least vulnerable agent across models. This is due to ExpeL's trajectory-based memory, which treats memories as prior examples rather than authoritative facts, reducing the likelihood of utilizing malicious or misleading information from poisoned memories. The only

other exception is MemoryOS on Llama, with an activation rate of 42%. We observed that this agent and model combination struggled to even execute the requested task, as it would request more information or make incorrect tool calls, which naturally led to a lower attack activation rate. Overall, we see that the attack success rate is highest for ChatGPT and DeepSeek. Gemini and Llama are the least vulnerable to these attacks, but the attack success rate is still significant, with the lowest being ExpeL paired with Llama at 28%. Overall, the Llama model has the lowest attack activation rates, while ChatGPT and DeepSeek have the highest.

We analyzed attack activation across each attack category for both directive and descriptive payloads, as shown in Figures 6a and 6b. The heatmaps show the activation



图 6: *GhostWriter*的成功率结果，按内存注入（P1）成功率、攻击检索成功率及攻击激活成功率（P2）进行细分。




图 7: 按类别划分、并按模型与智能体分组的 P2 激活率 (%)。

在各类模型中，Mem0 和 ExpeL 的攻击检索率通常低于其他智能体；其中 Gemini 和 Llama 的检索率分别低至 57% 和 65%。这是因为 Gemini 和 Llama 在将输入转换为记忆时准确度较低，从而降低了将该攻击载荷与目标良性提示语相匹配的可能性。

所有模型中的攻击激活率均较高，所有智能体与大语言模型（LLM）的平均激活率为 60%。ExpeL 的激活率最低，在所有 LLM 中介于 28% 至 63% 之间，使其成为在各模型中抗攻击能力最强的智能体。这得益于 ExpeL 采用的基于轨迹的记忆机制——该机制将记忆视为先验示例而非权威事实，从而降低了利用受污染记忆中恶意或误导性信息的可能性。另一个例外是基于 Llama 的 Memo-

ryOS，其激活率为 42%。我们观察到，该智能体与模型组合甚至难以完成所请求的任务——它会要求提供更多信息或执行错误的工具调用，这自然导致了较低的攻击激活率。总体而言，ChatGPT 和 DeepSeek 的攻击成功率最高；Gemini 与 Llama 对此类攻击的抗性最强，但其攻击成功率仍较高，其中 ExpeL 与 Llama 的组合攻击成功率最低，为 28%。总体来看，Llama 模型的攻击激活率最低。

各项比率，其中 ChatGPT 和 DeepSeek 的比率最高。

我们分析了指令型与描述型有效载荷在各攻击类别中的攻击激活情况，如图6a和6b所示。热力图展示了四种大语




Figure 7: Prompt injection results against all of our evaluated agents. Notably, ExpeL is very vulnerable to this attack.

rates across four LLMs paired with various agents. With directive payloads (Figure 6a), we see that integrity corruption, and covert exfiltration exhibit the highest activation rates, peaking at 100% for ChatGPT, DeepSeek, and Llama. Conversely, sensitive info leakage shows the lowest overall susceptibility across all LLM models. Interestingly, for descriptive payloads (Figure 6b), we see that sensitive info leakage performs actually rates comparable to integrity corruption and covert exfiltration. Scope escalation activates at much lower rates for descriptive payloads, as this attack requires ignoring the normal bounds set by the user, such as never sending emails without review, and the attack payloads are not presented as strict instructions; the agents are more likely to treat them as optional guidelines to be ignored. In general, we also see the attack activation rates for descriptive payloads uniformly fall for integrity corruption and covert exfiltration with Llama, whereas in other models they are more varied, with some rising and others falling.

7.3. Prompt Injection Attacks and Defenses.

We also evaluated the impact of prompt-injection attacks on these long-term memory agents. We reused the experimental setup for *GhostWriter*, using a prompt-injection payload from the AgentDojo benchmark [12]. Our results are presented in Figure 7, which shows that prompt injection fails almost 100% of the time against most of these agents, except for ExpeL and MemoryOS. Prompt injection had activation rates of 1-22% against MemoryOS and 92% against ExpeL. ExpeL’s trajectory-based design causes it to almost always follow the directives from the prompt injection attack due to their authoritative nature. Additionally, we found that while P1 injections rate were very similar to *GhostWriter*, the P2 activation rates were very low, ranging from 0% to 16.7%. This is due to low retrieval rates and the payloads being nonoptimized for activation in later prompts.

We also wanted to determine if prompt injection detection methods would filter out *GhostWriter* payloads. To this end, we ran our directive and descriptive payloads through two prompt-injection detection methods to evaluate whether *GhostWriter* can be detected by them. We used

DataFilter [37], a Llama model trained to detect prompt-injection attacks, and PromptArmor [38], a prompt-only defense where we used ChatGPT. We found that DataFilter had a 0% detection rate against all of our attack payloads. Surprisingly, it was more effective to ask ChatGPT to look for prompt injection, as it actually had an 85% detection rate against our directive payloads, but that fell to only 6% against our descriptive payloads. We noticed that ChatGPT tends to classify payloads as prompt injection when they use authoritative language, which is why it can reliably detect our directive payloads. However, our descriptive payloads use more polite language, easily fooling this method, whereas *AM-Sentry* can reliably detect the attack regardless of whether the payloads are directive or descriptive.

7.4. AM-Sentry Effectiveness

Experiment Design: *AM-Sentry* has defenses at ingestion and at retrieval, so we separated the test for each phase. For the first-phase test, we compared the results of S_1 , S_2 , and S_3 to assess their effectiveness. We also compared against A-MAC [9], using its default parameters. A-MAC is not focused on vulnerability detection, and any attack payload it blocks is incidental, as its mechanisms only evaluate the utility of saving new memories. For the second-phase test, we combine the retrieval screen with S_1 , S_2 , and S_3 , and compare the overall results to determine the combined effectiveness of our approach. We began by generating a new snapshot for each *AM-Sentry* policy configuration and reused our adversary setup from Section 7.2. We evaluated our defenses against both directive and descriptive payloads, but for brevity, we discuss only our results against descriptive payloads, which were the hardest to detect.

Result Analysis: Figure 8 shows the effectiveness of our memory saving policies at preventing malicious memory injection. We first evaluate the A-MAC baseline, which was optimized solely for utility and does not consider vulnerability. Compared to our proposed policies, A-MAC exhibits greater variability across agents and inadvertently permits most attacks. The minimum attack success rate observed under A-MAC was 54% with an average of 84%. The lowest admission rates were under DeepSeek, which tended to score our descriptive payloads very low on utility.

S_1 admits almost everything under Gemini, but for DeepSeek and ChatGPT, it admits around 80% and 97% respectively. In contrast, Llama tends to indiscriminately overblock, considering the fact that S_1 only evaluates if a given memory will be useful or not. As a result, Llama is prone to filtering both malicious and legitimate payloads, which negatively impacts the agent’s operations.

S_2 reduces the admission of poisoned memories across all models except Llama. Notably, while Llama exhibits a lower block rate under S_2 than under S_1 , it’s less prone to blocking legitimate payloads, thereby improving the agent’s overall utility (detailed analysis in Section 7.5). ChatGPT benefits the most from S_2 , bringing the attack success rate down to 71% at the lowest. Gemini also benefits significantly from S_2 ’s improved structure, reducing attack success




图 7: 针对我们所有评估的智能体所获得的 Prompt Injection 攻击结果。值得注意的是, ExpeL 对这种攻击尤为脆弱。

言模型 (LLM) 与不同智能体配合时的攻击激活率。对于指令型有效载荷 (图6a), 我们发现完整性破坏和隐蔽数据外泄的攻击激活率最高, 在 ChatGPT、DeepSeek 和 Llama 上均达到 100%; 相反, 敏感信息泄露在所有 LLM 模型中整体受攻击性最低。有趣的是, 对于描述型有效载荷 (图6b), 敏感信息泄露的攻击激活率实际上与完整性破坏及隐蔽数据外泄相当。对于描述型有效载荷, 范围提升攻击的激活率要低得多, 因为此类攻击需要用户忽略其设定的常规行为边界 (例如: 绝不未经审核就发送电子邮件), 且攻击有效载荷并非以严格的指令形式呈现; 因此, 智能体更有可能将这些有效载荷视为可忽略的可选指南。总体而言, 我们还发现: 在使用 Llama 模型时, 描述性有效载荷的攻击激活率在完整性破坏和隐蔽数据外泄场景下均呈均匀下降趋势; 而在其他模型中, 该激活率则变化较大——部分模型的激活率上升, 而另一些模型的激活率则下降。

7.3. 即时注入攻击与防御措施。

我们还评估了提示词注入攻击对这些长时记忆智能体的影响。我们复用 GhostWriter 的实验设置, 采用来自 AgentDojo 基准测试集[12]的提示词注入有效载荷。实验结果如图7所示: 对于大多数智能体而言, 提示词注入攻击的失败率接近 100%, 但 ExpeL 和 MemoryOS 除外; 针对 MemoryOS, 提示词注入的激活率为 1–22%; 而针对 ExpeL, 激活率高达 92%。ExpeL 采用基于轨迹的设计方案, 由于其指令具有权威性, 因此几乎总是会遵循提示词注入攻击所发出的指令。此外, 我们发现尽管 P1 的注入成功率与 GhostWriter 非常接近, 但 P2 的激活率极低, 范围在 0% 至 16.7% 之间。这归因于较低的检索率以及该有效载荷未针对后续提示词中的激活操作进行优化。

我们还希望确定即时注入检测方法是否能够过滤掉 GhostWriter 的恶意载荷。为此, 我们将指令型和描述型恶意载荷分别输入到两种提示注入检测方法中, 以评估 GhostWriter 是否能被这些方法检测到。我们采用了 Data-

Filter[37] (一个经过训练用于检测提示注入攻击的 Llama 模型) 以及 PromptArmor[38] (一种仅基于提示的防御机制, 其中我们使用了 ChatGPT)。结果发现, DataFilter 对所有测试攻击载荷的检测率均为 0%。出乎意料的是, 让 ChatGPT 进行提示注入检测效果更佳: 其对指令型载荷的检测率高达 85%, 但对描述型载荷的检测率仅为 6%。我们注意到, 当恶意载荷使用权威性语言时, ChatGPT 往往会将其判定为提示注入攻击, 这正是其能够可靠检测到指令型载荷的原因; 然而, 我们的描述型载荷采用更为礼貌的语言, 极易欺骗该检测方法; 相比之下, AM-Sentry 无论恶意载荷属于指令型还是描述型, 都能可靠地检测出此类攻击。

7.4. AM-Sentry 效能

实验设计: AM-Sentry 在内存摄入阶段和内存检索阶段均设有防御机制, 因此我们将各阶段的测试分开进行。在第一阶段测试中, 我们对比了 S_1 , S_2 与 S_3 的测试结果以评估其有效性; 同时, 我们使用 A-MAC[9] 的默认参数对其进行对比。A-MAC 并非专注于漏洞检测, 其拦截的任何攻击载荷均为附带性结果, 因为其机制仅评估保存新内存的效用。在第二阶段测试中, 我们将内存检索筛选机制与 S_1 , S_2 及 S_3 相结合, 并对比整体结果以确定本方法的有效性。我们首先为每种 AM-Sentry 策略配置生成新的快照, 并复用第7.2节中设置的攻击者环境。我们评估了本系统对指令型与描述型载荷的防御效果, 但为简洁起见, 本文仅讨论针对描述型载荷的测试结果——这类载荷最难被检测到。

结果分析: 图8展示了我们提出的内存节省策略在防范恶意内存注入方面的有效性。我们首先对 A-MAC 基线模型进行了评估——该模型仅针对效用进行了优化, 未考虑漏洞问题。与我们提出的策略相比, A-MAC 在不同智能体之间表现出更大的差异性, 且无意中允许了大多数攻击发生。在 A-MAC 模型下观察到的最低攻击成功率仅为 54%, 而平均值为 84%。最低的准入率出现在 DeepSeek 模型中, 该模型对我们的描述性载荷在效用评分上往往给出极低的分值。

S_1 在 Gemini 环境下几乎会允许所有内容通过; 然而对于 DeepSeek 和 ChatGPT 来说, 其允许通过率分别约为 80% 和 97%。相比之下, Llama 往往会不加区分地实施过度拦截——考虑到 S_1 仅评估给定记忆内容是否具有用处这一事实。因此, Llama 容易对恶意及合法的负载均进行过滤, 这会对智能体的运行产生负面影响。

S_2 降低了除 Llama 之外所有模型对中毒记忆的接纳率。值得注意的是, 尽管 Llama 在 S_2 方案下的阻塞率低于在 S_1 方案下的阻塞率, 但它对合法有效载荷的阻塞倾向较低, 从而提升了智能体的整体效用 (详细分析见第7.5节)。ChatGPT 从 S_2 方案中获益最大, 其攻击成功率最低降至 71%; Gemini 也从 S_2 方案优化后的结构中显著受益, 降低了攻击成功率。




Figure 8: Memory Saving policy performance; shows the number of malicious memories admitted by each policy into the agents’ memory store. Overall, S_3 performs best at blocking poisoned memory.




Figure 9: The total attack success rate for each policy configuration when paired with R. An attack is successful if it activates and alters the agent’s behavior. Pairing R with S_1 - S_3 greatly improves their overall performance against *GhostWriter*.

to between 70-84%. Conversely, DeepSeek is largely the same, and in some cases worse, because the model tends to rank the utility of our attack email fairly low. S_2 is consistent across most agents, but performance degrades with A-Mem and MemoryOS due to their memory structures, making S_2 more permissive in those agents and reducing its effectiveness.

S_3 is the most effective and consistent approach for most models. We observe improved performance across ChatGPT, DeepSeek, and Gemini, reducing the average attack success rate to 15%. A glaring exception is Llama, which seemed to struggle with the more strict, precise format of S_3 admitting 77% of malicious payloads. Overall, S_3 is highly effective in blocking the attack, provided it is paired with a capable LLM model.

We analyzed the total attack success rates when each policy configuration (S_1 - S_3) is integrated with the retrieval screening (R). Figure 9 shows the end-to-end attack success rates under the various policies. For an attack to be considered successful, it must pass both the policy and the retrieval screen. In this configuration, S_1 and S_2 exhibit comparable performance. The addition of the retrieval screen effectively mitigates the inherent permissiveness of S_1 and S_2 , bringing their performance to a similar level. S_3 is highly effective

in this combined setup, reducing the average attack success rate to below 12% across all LLMs except Llama, which averages 20% across all agents. We can see that R was able to make up for the weakness of S_3 exhibited under Llama during memory injection.

7.5. AM-Sentry Utility Analysis

Experiment Design: In addition to evaluating the mitigation’s effectiveness, we also evaluate the impact *AM-Sentry* has on each agent’s utility. This allows us to better understand the trade-off between protection and utility. For each agent across all models, we ran our test suite on the simulated work week to evaluate the agent’s recall and tool-calling accuracy. We repeated this evaluation for each *AM-Sentry* configuration and for A-MAC as an additional baseline. Each configuration created a snapshot based on the simulated work week and its memory-saving policies. In this subsection, we focus on A-Mem and present per-agent utility for the other agents in Appendix D.

Result Analysis: The utility results for A-Mem are presented in Figure 10, which gives the average utility scores for each configuration of *AM-Sentry*, A-MAC, and the baseline (no policy). We evaluated three metrics: F1-similarity




图 8：内存节省策略性能；展示了各策略允许被添加到智能体内存存储中的恶意内存条目数量。总体而言， S_3 在阻断受污染内存方面表现最佳。




图 9：各策略配置与 R 配合时的总攻击成功率。当攻击成功激活并改变了智能体的行为时，该攻击即视为成功。将 R 与 S_1 - S_3 配合使用可显著提升其对抗 *GhostWriter* 时的整体性能。

介于 70–84% 之间。相反，DeepSeek 的表现大体相当，甚至在某些情况下更差，这是因为该模型往往将我们发起的攻击性电子邮件的效用评分定得相对较低。 S_2 在大多数智能体中表现一致，但当面对 A-Mem 和 MemoryOS 时，其性能会因这两种智能体的内存结构而下降；这使得 S_2 在这些智能体中更具包容性，从而降低了其有效性。

S_3 对大多数模型而言是最有效且最稳定的防御方法。我们观察到 ChatGPT、DeepSeek 和 Gemini 的防御性能均有所提升，平均攻击成功率降至 15%。一个显著的例外是 Llama——该模型在面对 S_3 这种更严格、更精确的格式时似乎难以应对，其防御率仍高达 77%。总体而言，只要与性能优异的大型语言模型（LLM）结合使用， S_3 在阻断此类攻击方面效果显著。

我们分析了当每种策略配置（ S_1 - S_3 ）与检索筛选机制（R）相结合时的总体攻击成功率。图9展示了在不同策略下的端到端攻击成功率。要判定一次攻击为成功，该攻击必须同时通过策略筛选和检索筛选。在该配置下， S_1 与 S_2 的表现相当。引入检索筛选机制有效缓解了 S_1 和 S_2 固有的宽松性，使其性能趋于一致。 S_3 在这一组

合配置中效果显著，将所有大型语言模型（除Llama外）的平均攻击成功率降低至12%以下；而Llama在所有智能体中的平均攻击成功率为20%。由此可见，R能够弥补 S_3 在Llama进行记忆注入时所表现出的不足。

7.5. AM-Sentry实用分析

实验设计：除评估缓解措施的有效性外，我们还评估了 *AMSentry* 对各智能体效用的影响。这有助于我们更好地理解保护机制与效用之间的权衡关系。针对所有模型中的每个智能体，我们在模拟的工作周内运行了测试套件，以评估该智能体的召回率和工具调用准确率。我们对每种 *AM-Sentry* 配置以及作为附加基线的 A-MAC 进行了上述评估。每个配置均基于模拟的工作周及其节能策略创建了一个快照。在本小节中，我们重点讨论 A-Mem，并在附录D中呈现其余智能体的效用值。

结果分析：A-Mem 的效用结果如图10所示，该图展示了 *AM-Sentry*、A-MAC 以及基线（无策略）各配置方案的平均效用得分。我们评估了三项指标：F1 相似度




Figure 10: Utility metrics for A-Mem showing the F1-similarity score, LLM judged scores, and tool accuracy across the baseline, A-MAC, and all six *AM-Sentry* configurations. Overall, *AM-Sentry* has minimal impact on utility.

score, LLM-judge score, and tool accuracy. The F1-score measures how closely the generated answer matches a reference; the LLM-judge score uses a ChatGPT model to assess answer accuracy; and tool accuracy measures how often the correct tool call was made.

F1-scores remain consistent with the baseline across all LLMs, DeepSeek exhibiting the most consistent performance. A small reduction under S_3 is observed compared to the baseline, but we determined that this is an acceptable trade-off for protection. The remaining models show degraded performance when paired with R or S_3 , ranging from 0.01 to 0.04 relative to the baseline. With Llama, we also observe a stepwise reduction in utility starting at S_1 , due to indiscriminate over-blocking of legitimate memories. A-MAC is a utility-focused approach, but it did not have a strong positive effect on baseline performance in our experiments; the only exception is Llama, which scores significantly higher than the other policy configurations. We attribute this to A-MAC’s policy decisions being less reliant on LLMs, whereas our approach is more reliant on LLMs to score all parameters.

The LLM judged scores follow the same trend. Notably, S_2 outperforms the baseline in several cases. Tool accuracy is likewise consistent, with one exception. For ChatGPT, we observe a drop-off under A-MAC and S_1 relative to the baseline, and accuracy generally increases as the policy configuration becomes more strict. Notably, our policy configurations have minimal impact on overall utility, a trend that holds across other agent architectures as well, as further detailed in Appendix D.

8. Limitations & Future Work

In this paper, we focused on email and calendar invites as the attack surface. Other input channels, including documents, web content, and code repositories, were not explored. As an initial step in investigating this attack vector, our proposed defense relies on intuitively chosen weights in its equations, which may not be the optimal configuration for protecting against *GhostWriter*. Our utility benchmark is a custom test suite built around a simulated workweek rather than real user data or live deployments, and it may

not capture the full range of tasks a personal assistant agent encounters in practice.

Our evaluation covers five state-of-the-art memory agents, but the scope of memory architectures is broader. Agents with fundamentally different memory designs may exhibit different vulnerability profiles. We evaluated each agent using multiple models, but we were limited to those available and economically viable for us. We evaluated *AM-Sentry* against non-adaptive attackers unaware of the defense. An adaptive attacker may employ techniques to circumvent the policy checklist and retrieval screen.

From these limitations, we identify the following items for future work. Investigating *GhostWriter* across additional input channels and deployment contexts beyond the office assistant setting. Evaluating a wider range of memory architectures and additional LLM backends. Developing a dedicated benchmark spanning diverse task types, input sources, and memory architectures. Evaluating *AM-Sentry* against adaptive attackers, performing a Monte Carlo for optimizing selection of policy weights, and attack analysis with multiple concurrent attack payloads.

9. Conclusion

In this work, we introduced the novel attack vector *GhostWriter* which exploits memory systems in tool-using personal agents and poisons their memory store to steer their future behavior. We showed that it has near-universal success against current memory architectures and demonstrates high retrieval and activation rates due to the lack of security-focused memory governance. In response to these findings, we also proposed *AM-Sentry*, which utilized two mitigation techniques: a memory-saving policy and a retrieval screen. We show that *AM-Sentry* can greatly reduce the effectiveness of *GhostWriter*. The strictest version of *AM-Sentry* almost completely blocks *GhostWriter* across most models in our experiments. This results in a minimal drop in utility, introducing an acceptable trade-off between security and utility.

References

- [1] T. Schick, J. Dwivedi-Yu, R. Dessi, R. Raileanu, M. Lomeli, E. Hambro, L. Zettlemoyer, N. Cancedda, and T. Scialom, “Toolformer:




图 10: A-Mem 的实用度指标, 展示了基线模型、A-MAC 以及所有六种 *AM-Sentry* 配置下的 F1 相似度得分、LLM 评估得分及工具准确率。总体而言, *AM-Sentry* 对实用度的影响极小。

得分、LLM 评估分数及工具准确率。F1 分数用于衡量生成的回答与参考答案的匹配度; LLM 评估分数采用 ChatGPT 模型来评估回答的准确性; 而工具准确率则用于衡量正确调用工具的频率。

在所有大型语言模型 (LLM) 中, F1 分数与基线模型保持一致, 其中 DeepSeek 的性能最为稳定。与基线模型相比, 在 S_3 模型下观察到小幅下降, 但我们认定这一折衷方案对于保护目的而言是可接受的。其余模型在搭配 R 或 S_3 模型时, 其性能均有所下降, 下降幅度相对于基线模型介于 0.01 至 0.04 之间。对于 Llama 模型, 我们还观察到其效用值从 S_1 模型起开始逐步下降, 这是由于对合法记忆进行了无差别过度屏蔽所致。A-MAC 是一种以效用为核心的策略, 但在我们的实验中, 它并未对基线模型的性能产生显著的正向影响; 唯一的例外是 Llama 模型, 其得分显著高于其他策略配置。我们将此归因于 A-MAC 的策略决策对 LLM 的依赖程度较低, 而我们的方法则更依赖于 LLM 来对所有参数进行评分。

LLM 评判得分也呈现相同的趋势。值得注意的是, S_2 在多个案例中均优于基线模型。工具准确率同样保持稳定, 仅有一个例外: 对于 ChatGPT, 我们观察到在 A-MAC 和 S_1 模型中, 其准确率相较于基线模型有所下降; 而随着策略配置日趋严格, 准确率通常会提高。值得注意的是, 我们的策略配置对整体效用影响甚微——这一趋势在其他智能体架构中也同样成立, 具体细节详见附录 D。

8. 局限性与未来工作

在本文中, 我们重点研究电子邮件和日历邀请作为攻击面; 而其他输入渠道 (包括文档、网页内容及代码仓库) 则未被纳入研究范围。作为探索这一攻击向量的初步尝试, 我们提出的防御方案基于其方程中直观设定的权重参数, 但这些参数可能并非防御 *GhostWriter* 攻击的最佳配置。我们的实用性能基准测试采用了一套基于模拟工作周构建的定制化测试套件, 而非基于真实用户数

据或实际部署环境; 因此, 该测试套件可能无法全面反映个人助理代理在实际使用中所面临的各类任务场景。

我们的评估涵盖了五款最先进的记忆型智能体, 但记忆架构的范围更为广泛; 采用 *fundamentally* 不同记忆设计的智能体可能呈现出不同的漏洞特征谱。我们使用多种模型对每款智能体进行了评估, 但评估范围仅限于当前可获取且对我们的成本效益而言可行的模型。我们针对 *AM-Sentry* 进行了评估, 测试对象为对防御机制不知情的非自适应攻击者; 而自适应攻击者则可能采用特定技术来规避策略检查清单和检索筛选机制。

基于上述局限性, 我们为后续研究确定了以下工作方向: 在 Office 助理应用场景之外, 对 *GhostWriter* 进行更多输入通道和部署环境下的研究; 评估更广泛的内存架构及额外的大型语言模型 (LLM) 后端; 开发一个涵盖多样化任务类型、输入源及内存架构的专用基准测试集; 针对 *AM-Sentry* 进行对抗性攻击场景下的评估, 采用蒙特卡洛方法优化策略权重的选择, 并利用多个并发攻击载荷进行攻击分析。

9. 结论

在本研究中, 我们引入了一种新型攻击向量 *GhostWriter*, 该攻击向量利用工具使用型个人智能体中的内存系统, 并对它们的内存存储库进行“中毒”以操控其未来行为。我们证明了该攻击向量在当前各类内存架构上均具有近乎普遍的成功率, 并且由于缺乏以安全为核心的内存治理机制, 其检索与激活率较高。针对这些发现, 我们还提出了 *AM-Sentry*, 该方案采用了两种缓解技术: 内存节省策略与检索筛选机制。我们证明了 *AM-Sentry* 能显著降低 *GhostWriter* 的攻击效果。在我们的实验中, *AM-Sentry* 的最严格版本几乎完全阻断了 *GhostWriter* 在大多数模型上的攻击效果。这仅导致效用值出现微小下降, 从而在安全性与效用之间实现了可接受的权衡。

参考文献

- [1] T. Schick, J. Dwivedi-Yu, R. Dessì, R. Raileanu, M. Lomeli, E.

- Language models can teach themselves to use tools,” *Advances in neural information processing systems*, vol. 36, pp. 68 539–68 551, 2023.
- [2] L. Wang, C. Ma, X. Feng, Z. Zhang, H. Yang, J. Zhang, Z. Chen, J. Tang, X. Chen, Y. Lin *et al.*, “A survey on large language model based autonomous agents,” *Frontiers of Computer Science*, vol. 18, no. 6, p. 186345, 2024.
 - [3] C. Packer, V. Fang, S. Patil, K. Lin, S. Wooders, and J. Gonzalez, “MemGPT: towards LLMs as operating systems,” 2023.
 - [4] A. Zhao, D. Huang, Q. Xu, M. Lin, Y.-J. Liu, and G. Huang, “Expel: Llm agents are experiential learners,” in *Proceedings of the AAAI Conference on Artificial Intelligence*, vol. 38, no. 17, 2024, pp. 19 632–19 642.
 - [5] J. Kang, M. Ji, Z. Zhao, and T. Bai, “Memory os of ai agent,” in *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing*, 2025, pp. 25 972–25 981.
 - [6] T. Kagaya, T. J. Yuan, Y. Lou, J. Karlekar, S. Pranata, A. Kinose, K. Oguri, F. Wick, and Y. You, “Rap: Retrieval-augmented planning with contextual memory for multimodal llm agents,” *arXiv preprint arXiv:2402.03610*, 2024.
 - [7] OpenAI, “GPT-4 Technical Report,” 2024. [Online]. Available: <https://arxiv.org/abs/2303.08774>
 - [8] S. Wu, “Introducing Devin, the first AI software engineer,” <https://cognition.ai/blog/introducing-devin>, March 2024, cognition AI Blog. Accessed: 2025.
 - [9] G. Zhang, W. Jiang, X. Wang, A. Behr, K. Zhao, J. Friedman, X. Chu, and A. Anoun, “Adaptive Memory Admission Control for LLM Agents,” *arXiv preprint arXiv:2603.04549*, 2026.
 - [10] Z. Chen, Z. Xiang, C. Xiao, D. Song, and B. Li, “Agentpoison: Red-teaming llm agents via poisoning memory or knowledge bases,” *Advances in Neural Information Processing Systems*, vol. 37, pp. 130 185–130 213, 2024.
 - [11] S. Dong, S. Xu, P. He, Y. Li, J. Tang, T. Liu, H. Liu, and Z. Xiang, “Memory Injection Attacks on LLM Agents via Query-Only Interaction,” *arXiv preprint arXiv:2503.03704*, 2025.
 - [12] E. Debenedetti, J. Zhang, M. Balunovic, L. Beurer-Kellner, M. Fischer, and F. Tramèr, “Agentdojo: A dynamic environment to evaluate prompt injection attacks and defenses for llm agents,” *Advances in Neural Information Processing Systems*, vol. 37, pp. 82 895–82 920, 2024.
 - [13] T. Xie, F. Zhou, Z. Cheng, P. Shi, L. Weng, Y. Liu, T. J. Hua, J. Zhao, Q. Liu, C. Liu *et al.*, “Openagents: An open platform for language agents in the wild,” *arXiv preprint arXiv:2310.10634*, 2023.
 - [14] Significant Gravitas, “AutoGPT,” <https://github.com/Significant-Gravitas/AutoGPT>, 2023, accessed: 2026.
 - [15] Y. Nakajima, “BabyAGI,” <https://github.com/yoheinakajima/babyagi>, 2023, accessed: 2026.
 - [16] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafraan, K. R. Narasimhan, and Y. Cao, “React: Synergizing reasoning and acting in language models,” in *The eleventh international conference on learning representations*, 2022.
 - [17] Anthropic, “Claude Model Card,” <https://www.anthropic.com/model-card>, 2023, accessed: 2026.
 - [18] Z. Xi, W. Chen, X. Guo, W. He, Y. Ding, B. Hong, M. Zhang, J. Wang, S. Jin, E. Zhou *et al.*, “The rise and potential of large language model based agents: A survey,” *Science China Information Sciences*, vol. 68, no. 2, p. 121101, 2025.
 - [19] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel *et al.*, “Retrieval-augmented generation for knowledge-intensive nlp tasks,” *Advances in neural information processing systems*, vol. 33, pp. 9459–9474, 2020.
 - [20] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, Y. Dai, J. Sun, H. Wang, H. Wang *et al.*, “Retrieval-augmented generation for large language models: A survey,” *arXiv preprint arXiv:2312.10997*, vol. 2, no. 1, p. 32, 2023.
 - [21] J. S. Park, J. O’Brien, C. J. Cai, M. R. Morris, P. Liang, and M. S. Bernstein, “Generative agents: Interactive simulacra of human behavior,” in *Proceedings of the 36th annual acm symposium on user interface software and technology*, 2023, pp. 1–22.
 - [22] W. Xu, Z. Liang, K. Mei, H. Gao, J. Tan, and Y. Zhang, “A-mem: Agentic memory for llm agents,” *arXiv preprint arXiv:2502.12110*, 2025.
 - [23] N. Shinn, F. Cassano, E. Berman, A. Gopinath, K. Narasimhan, and S. Yao, “Reflexion: Language agents with verbal reinforcement learning, 2023,” *URL https://arxiv.org/abs/2303.11366*, vol. 8, 2024.
 - [24] L. Franceschi-Bicchierai, “Hackers hijacked instagram accounts by tricking meta ai support chatbot into granting access,” Jun. 2026, accessed: 2026-06-05. [Online]. Available: <https://techcrunch.com/2026/06/01/hackers-hijacked-instagram-accounts-by-tricking-meta-ai-support-chatbot-into-granting-access/>
 - [25] X. Yang, Y. He, S. Ji, B. Hooi, and J. S. Dong, “Zombie agents: Persistent control of self-evolving llm agents via self-reinforcing injections,” *arXiv preprint arXiv:2602.15654*, 2026.
 - [26] S. S. Srivastava and H. He, “MemoryGraft: Persistent compromise of LLM agents via poisoned experience retrieval,” *arXiv preprint arXiv:2512.16962*, 2025.
 - [27] Q. Wei, T. Yang, Y. Wang, X. Li, L. Li, Z. Yin, Y. Zhan, T. Holz, Z. Lin, and X. Wang, “A-memguard: A proactive defense framework for llm-based agent memory,” *arXiv preprint arXiv:2510.02373*, 2025.
 - [28] B. Klimt and Y. Yang, “The enron corpus: A new dataset for email classification research,” in *European conference on machine learning*. Springer, 2004, pp. 217–226.
 - [29] N. Reimers and I. Gurevych, “Sentence-bert: Sentence embeddings using siamese bert-networks,” 2019. [Online]. Available: <https://arxiv.org/abs/1908.10084>
 - [30] S. Xiao, Z. Liu, P. Zhang, N. Muennighoff, D. Lian, and J.-Y. Nie, “C-pack: Packed resources for general chinese embeddings,” 2024. [Online]. Available: <https://arxiv.org/abs/2309.07597>
 - [31] FBI Internet Crime Complaint Center (IC3), “Internet Crime Report 2023,” 2024, accessed: 2026-04-26. [Online]. Available: https://www.ic3.gov/Media/PDF/AnnualReport/2023_IC3Report.pdf
 - [32] P. Chhikara, D. Khant, S. Aryan, T. Singh, and D. Yadav, “Mem0: Building production-ready ai agents with scalable long-term memory,” *arXiv preprint arXiv:2504.19413*, 2025.
 - [33] Letta Team, “Letta: Build and deploy stateful agents,” <https://github.com/letta-ai/letta>, 2024, formerly MemGPT.
 - [34] A. Maharana, D.-H. Lee, S. Tulyakov, M. Bansal, F. Barbieri, and Y. Fang, “Evaluating very long-term conversational memory of llm agents,” in *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2024, pp. 13 851–13 870.
 - [35] Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. Cohen, R. Salakhutdinov, and C. D. Manning, “HotpotQA: A dataset for diverse, explainable multi-hop question answering,” in *Proceedings of the 2018 conference on empirical methods in natural language processing*, 2018, pp. 2369–2380.
 - [36] S. Yao, H. Chen, J. Yang, and K. Narasimhan, “Webshop: Towards scalable real-world web interaction with grounded language agents,” *Advances in Neural Information Processing Systems*, vol. 35, pp. 20 744–20 757, 2022.
 - [37] Y. Wang, S. Chen, R. Alkhudair, B. Alomair, and D. Wagner, “Defending against prompt injection with datafilter,” *arXiv preprint arXiv:2510.19207*, 2025.
 - [38] T. Shi, K. Zhu, Z. Wang, Y. Jia, W. Cai, W. Liang, H. Wang, H. Alzahrani, J. Lu, K. Kawaguchi *et al.*, “Promptarmor: Simple yet effective prompt injection defenses,” *arXiv preprint arXiv:2507.15219*, 2025.

- Hambro, L. Zettlemoyer, N. Cancedda 与 T. Scialom, “Toolformer: 语言模型可自主学习如何使用工具”, 《神经信息处理系统进展》, 第 36 卷, 第 68 539–68 551 页, 2023 年。
- [2] L. Wang, C. Ma, X. Feng, Z. Zhang, H. Yang, J. Zhang, Z. Chen, J. Tang, X. Chen, Y. Lin 等, 《基于大型语言模型的自主智能体综述》, 《计算机科学前沿》, 第 18 卷, 第 6 期, 第 186345 页, 2024 年。
- [3] C. Packer, V. Fang, S. Patil, K. Lin, S. Wooders 与 J. Gonzalez, 《MemGPT: 迈向将大型语言模型作为操作系统》, 2023 年。
- [4] A. Zhao, D. Huang, Q. Xu, M. Lin, Y.-J. Liu 与 G. Huang, 《Expel: LLM 代理是经验型学习者》, 载于《AAAI 人工智能会议论文集》, 第 38 卷, 第 17 期, 2024 年, 第 19 632–19 642 页。
- [5] J. Kang, M. Ji, Z. Zhao 与 T. Bai, 《AI 代理的内存》, 载于《2025 年自然语言处理实证方法会议论文集》, 2025 年, 第 25 972–25 981 页。
- [6] T. Kagaya, T. J. Yuan, Y. Lou, J. Karlekar, S. Pranata, A. Kinose, K. Oguri, F. Wick 与 Y. You, “Rap: 基于上下文记忆的检索增强规划方法用于多模态大语言模型智能体”, *arXiv 预印本 arXiv:2402.03610*, 2024 年。
- [7] OpenAI, 《GPT-4 Technical Report》, 2024 年。[在线]。访问地址: <https://arxiv.org/abs/2303.08774>
- [8] S. Wu, 《介绍 Devin——首位 AI 软件工程师》, <https://cognition.ai/blog/introducing-devin>, 2024 年 3 月, Cognition AI 博客。访问日期: 2025 年。
- [9] G. Zhang, W. Jiang, X. Wang, A. Behr, K. Zhao, J. Friedman, X. Chu 与 A. Anoun, “面向 LLM 智能体的自适应内存准入控制机制”, *arXiv 预印本 arXiv:2603.04549*, 2026 年。
- [10] Z. Chen, Z. Xiang, C. Xiao, D. Song 与 B. Li, 《Agentpoison: 通过污染记忆或知识库对人工智能代理进行红队攻击》, 《神经信息处理系统进展》, 第 37 卷, 第 130 185–130 213 页, 2024 年。
- [11] S. Dong, S. Xu, P. He, Y. Li, J. Tang, T. Liu, H. Liu 与 Z. Xiang, 《通过纯查询交互对大型语言模型智能体实施内存注入攻击》, *arXiv 预印本 arXiv:2503.03704*, 2025 年。
- [12] E. Debenedetti, J. Zhang, M. Balunovic, L. Beurer-Kellner, M. Fischer 与 F. Tramèr, “Agentdojo: 一种用于评估大语言模型智能体中提示注入攻击与防御机制的动态环境”, 《神经信息处理系统进展》, 第 37 卷, 第 8285–82 920 页, 2024 年。
- [13] 谢涛、周峰、程志、史鹏、翁乐、刘阳、华涛、赵杰、刘强、刘畅等, 《Openagents: 面向真实场景的语言智能体开放平台》, *arXiv 预印本 arXiv:2310.10634*, 2023 年。
- [14] Significant Gravitas, “AutoGPT”, <https://github.com/Significant-Gravitas/AutoGPT>, 2023 年, 访问日期: 2026 年。
- [15] Y. Nakajima, 《BabyAGI》, <https://github.com/yoheinakajima/babyagi>, 2023 年, 访问日期: 2026 年。
- [16] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafraan, K. R. Narasimhan 与 Y. Cao, “React: 在语言模型中实现推理与行动的协同作用”, 载于《第十一届国际学习表示会议》, 2022 年。
- [17] Anthropic, 《Claude 模型卡片》, <https://www.anthropic.com/model-card>, 2023 年, 访问日期: 2026 年。
- [18] Z. Xi, W. Chen, X. Guo, W. He, Y. Ding, B. Hong, M. Zhang, J. Wang, S. Jin, E. Zhou 等, 《基于大型语言模型的智能体的发展与潜力: 综述》, 《中国科学: 信息科学》, 第 68 卷, 第 2 期, 第 121101 页, 2025 年。
- [19] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel 等, 《基于检索增强的生成方法在知识密集型自然语言处理任务中的应用》, 《神经信息处理系统进展》, 第 33 卷, 第 9459–9474 页, 2020 年。
- [20] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, Y. Dai, J. Sun, H. Wang, H. Wang 等, 《大型语言模型中的检索增强生成: 综述》, *arXiv 预印本 arXiv:2312.10997*, 第 2 卷, 第 1 期, 第 32 页, 2023 年。
- [21] J. S. Park, J. O. ‘Brien, C. J. Cai, M. R. Morris, P. Liang 与 M. S. Bernstein, 《生成式智能体: 人类行为的交互式拟像》, 载于《第 36 届 ACM 用户界面软件与技术年会论文集》, 2023 年, 第 1–22 页。
- [22] W. Xu, Z. Liang, K. Mei, H. Gao, J. Tan 与 Y. Zhang, 《A-mem: 面向大语言模型代理的代理式记忆》, *arXiv 预印本 arXiv:2502.12110*, 2025 年。
- [23] N. Shinn, F. Cassano, E. Berman, A. Gopinath, K. Narasimhan 与 S. Yao, “Reflexion: 基于语言强化学习的语言智能体, 2023”, [网址 https://arxiv.org/abs/2303.11366](https://arxiv.org/abs/2303.11366), 第 8 卷, 2024 年。
- [24] L. Franceschi-Bicchieri, 《黑客通过欺骗 Meta AI 支持聊天机器人获取访问权限而劫持 Instagram 账户》, 2026 年 6 月, 访问日期: 2026-06-05。[网络版]。获取地址: <https://techcrunch.com/2026/06/01/hackers-hijacked-instagram-accounts-by-tricking-meta-ai-support-chatbot-into-granting-access/>
- [25] X. Yang, Y. He, S. Ji, B. Hooi 与 J. S. Dong, “Zombie agents: 通过自强化注入实现自演化 LLM 代理的持续控制”, *arXiv 预印本 arXiv:2602.15654*, 2026 年。
- [26] S. S. Srivastava 与 H. He, “MemoryGraft: 通过中毒式经验检索对大型语言模型智能体进行持久性攻破”, *arXiv 预印本 arXiv:2512.16962*, 2025 年。
- [27] Q. Wei, T. Yang, Y. Wang, X. Li, L. Li, Z. Yin, Y. Zhan, T. Holz, Z. Lin 与 X. Wang, “A-memguard: 一种基于 LLM 的智能体内存的主动防御框架”, *arXiv 预印本 arXiv:2510.02373*, 2025 年。
- [28] B. Klimt 与 Y. Yang, 《Enron 语料库: 用于电子邮件分类研究的新数据集》, 载于《欧洲机器学习会议》, Springer, 2004 年, 第 217–226 页。
- [29] N. Reimers 与 I. Gurevych, 《Sentence-BERT: 基于孪生 BERT 网络的句子嵌入方法》, 2019 年。[在线]。获取地址: <https://arxiv.org/abs/1908.10084>
- [30] S. Xiao, Z. Liu, P. Zhang, N. Muennighoff, D. Lian 与 J.-Y. Nie, 《C-Pack: 适用于通用中文嵌入的打包资源》, 2024 年。[在线]。获取地址: <https://arxiv.org/abs/2309.07597>
- [31] 美国联邦调查局互联网犯罪投诉中心 (IC3) 《2023 年互联网犯罪报告》, 2024 年, 访问日期: 2026 年 4 月 26 日。[在线]。获取地址: https://www.ic3.gov/Media/PDF/AnnualReport/2023_IC3Report.pdf
- [32] P. Chhikara, D. Khant, S. Aryan, T. Singh 与 D. Yadav, “Mem0: 构建具备可扩展长期记忆功能的、可投入实际生产的人工智能代理”, *arXiv 预印本 arXiv:2504.19413*, 2025 年。
- [33] Letta 团队, 《Letta: 构建与部署有状态智能体》, <https://github.com/letta-ai/letta>, 2024 年; 原名 MemGPT。
- [34] A. Maharana, D.-H. Lee, S. Tulyakov, M. Bansal, F. Barbieri 与 Y. Fang, 《《评估大型语言模型智能体的超长期对话记忆》》, 载于《第 62 届计算语言学学会年会论文集 (第 1 卷: 长文)》, 2024 年, 第 13 851–13 870 页。
- [35] Z. Yang, P. Qi, S. Zhang, Y. Bengio, W. Cohen, R. Salakhutdinov 与 C. D. Manning, “HotpotQA: 一个用于多样化、可解释性多跳问答系统的数据集”, 载于《2018 年自然语言处理经验方法会议论文集》, 2018 年, 第 2369–2380 页。
- [36] S. Yao, H. Chen, J. Yang 与 K. Narasimhan, “Webshop: 面向基于现实场景的语言智能体的可扩展现实世界 Web 交互系统”, 《神经信息处理系统进展》, 第 35 卷, 第 20 744–20 757 页, 2022 年。
- [37] Y. Wang, S. Chen, R. Alkhudair, B. Alomair 与 D. Wagner, 《利用 DataFilter 防范即时注入攻击》, *arXiv 预印本 arXiv:2510.19207*, 2025 年。
- [38] T. Shi, K. Zhu, Z. Wang, Y. Jia, W. Cai, W. Liang, H. Wang, H. Alzahrani, J. Lu, K. Kawaguchi 等, 《PromptArmor: 简单而有效的提示注入防御机制》, *arXiv 预印本 arXiv:2507.15219*, 2025 年。




Figure 11: A directive attack email that uses explicit instructions and autocratic language.

Appendix A. Ethical Considerations

This paper studies security vulnerabilities in personal assistant agents that use long-term memory. We formalize and evaluate an attack setup against long-term-memory agents from the literature in a controlled environment using synthetic tasks and data; our experiments do not involve real users, sensitive user data, or third-party systems. Our evaluation is limited to author-controlled systems, and we focus on measuring the attack’s impact and developing defenses. The attack technique could be misused to compromise deployed agents with memory systems, so we also propose and evaluate a practical mitigation that substantially reduces the attack’s effectiveness. We plan to release our codebase to support reproducibility and further defensive research.

Appendix B. Adversary Emails

Figures 11 and 12 show an example of the same attack scenario across two payload variations, directive and descriptive. Directive payloads use more demanding and authoritative language, while descriptive payloads use more passive and polite language. While both are effective at driving agent behavior once stored in memory, their activation profiles differ slightly. The descriptive payloads are more often flagged as benign, while the directive payloads raise more red flags. All of our 16 scenarios and their variations will be included in our released codebase.




Figure 12: A descriptive attack email that introduces the malicious contact as a team member.

Appendix C. Synthetic Work Week

We present five examples of inputs from our simulated workweek in Figures 13 through 17. One email is about team coordination and scheduling a lunch, while the other is a technical discussion on a current issue. Our workweek consists of 32 events, combining both emails and calendar invites. All of these events are organized into a timeline of events and fed to the agents in chronological order. These inputs will be included in our released codebase.

Appendix D. Remaining AM-Sentry Utility Analysis

Figure 18 outlines the remaining utility results not included in the main paper. Overall, the trends align with those seen in A-Mem with a few notable exceptions. Llama tends to show much more weakness and variability for these agents. In particular, MemoryOS shows high variance when paired with Llama, while ExpeL and Letta degrade in performance as we switch between policy configurations. ExpeL, in general, also seems to suffer as policy configurations grow more strict.




图 11：一封采用明确指令和专制式语言的定向攻击电子邮件。

附录A。
伦理考量

本文研究了采用长期记忆机制的个人助理智能体中存在的安全漏洞。我们基于现有文献，利用合成任务和数据在受控环境中对基于长期记忆的智能体进行了攻击场景的形式化建模与评估；我们的实验不涉及真实用户、敏感用户数据或第三方系统。本次评估仅限于作者可控的系统，重点在于衡量攻击的影响程度并开发相应的防御机制。鉴于该攻击技术可能被恶意利用以攻破部署了记忆系统的智能体，我们还提出并评估了一种实用的缓解方案，可显著降低该攻击的有效性。我们计划公开相关代码库，以支持研究结果的可复现性及进一步的防御性研究。

附录B。
敌对邮件

图11与图12展示了同一攻击场景在两种负载变体（指令型与描述型）下的示例。指令型负载采用更具强制性和权威性的语言，而描述型负载则采用更为被动且礼貌的措辞。尽管这两种负载在存储于内存后均能有效驱动代理行为，但其激活特征存在细微差异：描述型负载更常被标记为“良性”，而指令型负载则会引发更多警觉。我们设计的全部 16 个场景及其变体均将包含在已发布的代码库中。




图 12：一封描述性攻击性电子邮件，将恶意联系人介绍为团队成员。

附录C。
合成工作周

我们在图13至17中展示了来自我们模拟工作周的五个工作输入示例。其中一个电子邮件涉及团队协调与午餐安排，另一个则是关于当前某个议题的技术讨论。我们的模拟工作周包含 32 个事件，涵盖了电子邮件和日历邀请两种形式。所有这些事件均被整理成事件时间线，并按时间顺序馈送至智能体。这些输入数据将包含在我们发布的代码库中。

附录 D。
剩余 AM-Sentry 实用工具分析

图18概述了未包含在主论文中的其余实用性能测试结果。总体而言，这些趋势与在 A-Mem 中观察到的趋势一致，但存在若干显著例外。对于这些智能体而言，Llama 的表现往往更为疲软且波动较大；具体而言，当 MemoryOS 与 Llama 配合使用时，其表现存在较高方差；而 ExpeL 和 Letta 的性能则会随着策略配置的切换而下降。总体来看，随着策略配置变得更为严格，ExpeL 的表现似乎也有所下降。




Figure 13: An internal email from a colleague coordinating a team lunch.




Figure 14: An email from a developer discussing implementation options for API rate limiting and reporting a pagination bug.




Figure 15: A calendar event for a team lunch.




Figure 16: A calendar update moving the date for a client demo.




Figure 17: An announcement notifying the change in location for the daily standup meeting.




图 13：一位同事发送的内部电子邮件，用于协调团队午餐事宜。




图 14：一封来自开发人员的电子邮件，其中讨论了 API 限速功能的实现方案，并报告了一个分页相关 Bug。




图 15：团队午餐的日历事件。




图 16：日历更新功能，用于调整客户演示会的日期。






图 17：用于通知每日站会地点变更的公告。



A: ChatGPT B: DeepSeek C: Gemini D: Llama

Figure 18: Utility metrics for Mem0, Letta, Expel, and MemoryOS. Showcasing the F1 similarity score, LLM judged scores, and tool accuracy across the baseline, A-MAC, and all six *AM-Sentry* configurations.



A: ChatGPT B: DeepSeek C: Gemini D: Llama

图 18: Mem0、Letta、Expel 和 MemoryOS 的实用性指标。展示了 F1 相似度得分、LLM 评估得分以及在基线模型、A-MAC 以及所有六种 *AM-Sentry* 配置下的工具准确率。